

Getting data out of Postgres: an update

Ian D. Gow

8 October 2026

Abstract The Arrow Parquet file format has emerged in recent years as a key enabler of several technologies that have the potential to revolutionize data analysis. For example, Polars is a “next-generation” data frame library that is much faster and more memory-efficient than pandas, the dominant data frame library for Python. However, the performance and efficiency of Polars depends on having the data in an efficient form such as is offered by the Parquet format. Polars really loses its edge when the data it wants are located in a relational database. In such a case, a natural question is: How do get data out of the database and into Parquet files?

This note explores a few ways of getting Parquet files from a PostgreSQL database, showing how to use them and comparing them with each other.

This is an update to my [26 March 2026 note](#), with new package versions and benchmarks using my local PostgreSQL database.

1 PostgreSQL

In 2011, I was looking for a place to store my data that didn’t require SAS, or zillions of Stata files. After dabbling with SQLite and MySQL, I landed on PostgreSQL as the best system for storing my data. PostgreSQL allowed me extremely reliable storage of large amounts of data with fast access from R, Python, Stata and more.

The challenge back then was: How do I get my data into PostgreSQL? At the time, I was a faculty member at Northwestern University and a lot of my data came from Wharton Research Data Services, more commonly referred to as WRDS (pronounced “words”), and at that time WRDS provided data as SAS data files.

To get data from SAS data files into PostgreSQL, I used SAS `PROC EXPORT` to create CSV files and the PostgreSQL commands to get the data from the CSV files into PostgreSQL. Over time, this approach morphed into Perl scripts, then Python scripts, and finally the Python package `wrds2pg`, which was [published on PyPI](#) in April 2020. This package essentially runs SAS code that emits CSV data that is piped directly in to PostgreSQL.

When I first started using PostgreSQL, I remember I often had to compile it from source and it was pretty complicated. Today, I can just install [Postgres.app](#) using a point-and-click package and I can have a PostgreSQL database running with the simplest defaults ready to receive data in minutes.¹

¹Unfortunately, Postgres.app is MacOS only.

2 What's the deal with Parquet?

One of the cool things about Parquet files is that any modern data science tool can work with them. Here I consider pandas, Polars and DuckDB. The examples use local copies of the data, and all database queries executed for this update connect to my local PostgreSQL database, [igow](#).

2.1 pandas

If you are a user of Python for data analysis, no doubt you do a lot of work with pandas, which has been the de facto standard data frame library in Python since about 2012.

```
import pandas as pd
```

Unlike systems such as Stata or SAS, there is no “pandas data file” format as such. My impression is that many pandas users store data as CSV files. Here I will focus on `crsp.msf_v2`, a data set that is at about the sweet spot in terms of size and complexity for illustrating some key ideas. For comparison, I also create a CSV copy of this data:

```
from pathlib import Path

parquet_file = Path("msf_v2.parquet")
csv_file = Path("msf_v2.csv")
pd.read_parquet(parquet_file).to_csv(csv_file, index=False, header=True)
```

If I have `crsp.msf_v2` as a CSV file, I can read it in using `pd.read_csv()`:

```
%%time
msf_v2 = pd.read_csv(csv_file, low_memory=False)
```

```
CPU times: user 16.8 s, sys: 1.71 s, total: 18.5 s
Wall time: 18.6 s
```

But if I had it as a Parquet file, I might be impressed with the performance of `pd.read_parquet()`:

```
%%time
msf_v2 = pd.read_parquet(parquet_file)
```

```
CPU times: user 8.63 s, sys: 7.01 s, total: 15.6 s
Wall time: 13.5 s
```

2.2 Polars

OK, I would only be impressed if I'd never used Polars, but I have used it and I import it here.²

```
import polars as pl
```

Now, how does `pl.read_parquet()` do with this file?

²When I wrote the original March note, I had only been using Polars for about two months.

```
%%time
msf_v2 = pl.read_parquet(parquet_file)
```

```
CPU times: user 1.11 s, sys: 2.43 s, total: 3.54 s
Wall time: 2.21 s
```

Again, if I didn't know Polars, I might be impressed. But in general you don't want to use something as *slow* as `pl.read_parquet()` for many tasks. Instead, just use `pl.scan_parquet()`.

```
%%time
msf_v2 = pl.scan_parquet(parquet_file)
```

```
CPU times: user 1.98 ms, sys: 16.4 ms, total: 18.4 ms
Wall time: 18.4 ms
```

OK, now I'm starting to be impressed. The surprising performance of `pl.scan_parquet()` comes from the fact that it's not actually reading all the data into memory. And this saves both time *and* memory.

Why is this helpful? Well, the reality is that we rarely want all the columns or all the rows, and the Parquet file format allows Polars to quickly pull in precisely the data that the analyst wants. Suppose we were interested in data from 2023 and just wanted returns and prices. This will do the trick:

```
%%time
msf_v2_2023 = (
    msf_v2
    .filter(pl.col("mthprcdt").dt.year() == 2023)
    .select("permno", "mthprcdt", "mthret", "mthprc")
    .collect()
)
```

```
CPU times: user 66.8 ms, sys: 31.5 ms, total: 98.3 ms
Wall time: 11.6 ms
```

And to prove that something happened, let's look at the data:

```
msf_v2_2023
```

permno	mthprcdt	mthret	mthprc
i32	date	decimal[10,6]	decimal[13,6]
10026	2023-01-31	-0.042816	143.300000
10026	2023-02-28	-0.014585	141.210000
10026	2023-03-31	0.054785	148.220000
10026	2023-04-28	0.033599	153.200000
10026	2023-05-31	0.004896	153.950000
...	...	...	...

permno	mothprcdt	mothret	mothprc
i32	date	decimal[10,6]	decimal[13,6]
93436	2023-08-31	-0.034962	258.080000
93436	2023-09-29	-0.030456	250.220000
93436	2023-10-31	-0.197346	200.840000
93436	2023-11-30	0.195379	240.080000
93436	2023-12-29	0.034988	248.480000

So, if `msf_v2_2023` were what we wanted, Polars gets it us hundreds of times faster than using pandas does, and while using much less RAM.³ These differences are probably not such a big deal with “medium-sized” tables such as `crsp.msf_v2`. But, if you’re working with larger data, it can quickly become simply impossible to work with data in pandas without using various complicating approaches, such as batching analyses and pasting together the results.⁴

2.3 DuckDB

Another tool that works well with Parquet files is DuckDB. DuckDB offers a rich SQL covering much of what PostgreSQL and in many areas even more. Unlike PostgreSQL, one does not need to set up a server to use DuckDB. I can simply import the library.

```
import duckdb
```

Then I can use the `duckdb.read_parquet()` function to do much the same thing as `pl.scan_parquet()`.⁵

```
%%time
msf_v2 = duckdb.read_parquet(parquet_file)
```

```
CPU times: user 4.07 ms, sys: 1.81 ms, total: 5.88 ms
Wall time: 5.7 ms
```

And much like with Polars I can then quickly get precisely the data I need. There are many ways to work with DuckDB and one interface uses SQL.⁶ The `.pl()` at the end returns a Polars data frame (`.df()` would give me a pandas data frame).

```
%%time
msf_v2_2023 = duckdb.sql("""
    SELECT permno, mothprcdt, mothret, mothprc
    FROM msf_v2
    WHERE year(mothprcdt) = 2023
""").pl()
```

```
CPU times: user 94.8 ms, sys: 19.5 ms, total: 114 ms
Wall time: 17.6 ms
```

³The pandas version of `msf_v2` occupies more than 8 GB of RAM.

⁴Another benefit is that the Parquet file occupies about 243 MB, versus 1.45 GB for the CSV file.

⁵Yes, this is a bit confusing.

⁶You can also use an SQL interface with Polars.

From the following, we can see that we’ve got the same data we got through the Polars path above:

```
msf_v2_2023
```

permno	nthprcdt	nthret	nthprc
i32	date	decimal[10,6]	decimal[13,6]
10026	2023-01-31	-0.042816	143.300000
10026	2023-02-28	-0.014585	141.210000
10026	2023-03-31	0.054785	148.220000
10026	2023-04-28	0.033599	153.200000
10026	2023-05-31	0.004896	153.950000
...	...	...	...
93436	2023-08-31	-0.034962	258.080000
93436	2023-09-29	-0.030456	250.220000
93436	2023-10-31	-0.197346	200.840000
93436	2023-11-30	0.195379	240.080000
93436	2023-12-29	0.034988	248.480000

3 Getting data into Parquet form

In late 2023, I decided to make a “Parquet path” for the book I was then working on *Empirical Research in Accounting: Tools and Methods*, which was later [published with CRC Press](#) and which used R for data analysis. This “Parquet path” combined Parquet files with DuckDB and the `dbplyr` library in R to create an alternative to the main “PostgreSQL path” of that book with only very minor tweaks to the code.

For this purposes, I pulled together a quick-and-dirty package `db2pq`, which creates Parquet files from PostgreSQL sources, including special functions to manage a local repository of WRDS data.⁷

Recently, I have noticed a lot of buzz around ADBC, which stands for Arrow Database Connectivity. The Apache Arrow project is closely related to the Apache Parquet project and I figured that that ADBC might provide a better way to get data from PostgreSQL to Parquet files.

In this note, I give some first impressions of using ADBC for this purpose and compare it with two alternative approaches.

3.1 Method 1: The DuckDB path of `db2pq`

The `db2pq` Python package allows a user to export a table from PostgreSQL to Parquet with a number of options, including the ability to select certain columns or filter only rows meaning certain conditions. It also allows the user to modify data types (e.g., have a floating-point number be read as an integer).

The main “engine” of `db2pq` has DuckDB read data from PostgreSQL and produce batches of Arrow output that are written to disk. The batching is necessary to ensure that large tables can be processed without swamping the memory of the user’s computer.

⁷For more on this, see [here](#).

While DuckDB supports writing directly to Parquet files, when I worked with this in 2023, I found it very difficult to prevent DuckDB from flooding RAM, and so I implemented a suite of functions that make up the core of the `db2pq` package.

The `db2pq` package also offers other goodies, such as only writing files when the data in PostgreSQL is newer than that in the Parquet file, as illustrated here:

i Note

The `wrds_update_pq()` examples in this section and the next retain their output from March 2026. I no longer have access to WRDS, so these examples are not rerun when rendering this note. The updated benchmarks below use my local PostgreSQL database.

```
from db2pq import wrds_update_pq
```

```
wrds_update_pq("msf_v2", "crsp")
```

```
Updated crsp.msf_v2 is available.  
Beginning file download at 2026-03-26 19:56:53 UTC.  
Completed file download at 2026-03-26 20:00:28 UTC.
```

```
wrds_update_pq("msf_v2", "crsp")
```

```
crsp.msf_v2 already up to date.
```

3.2 Method 2: The ADBC path of `db2pq`

In March 2026, I decided to add an ADBC “engine” to `db2pq`. I was interested to know whether it would deliver better performance. It turns out that it doesn’t—at least, not yet—but it was a useful exercise to understand a little more about ADBC.

To use the ADBC engine with `db2pq`, I simply say `engine="adbc"` when calling `wrds_update_pq()`.

```
wrds_update_pq("msf_v2", "crsp", engine="adbc")
```

```
crsp.msf_v2 already up to date.
```

Here I discuss two among several things I learnt about ADBC. First, the ADBC drivers incur a significant one-time cost when setting a database for connection. As we will see, the first export takes longer on my computer; in the original March 2026 tests, setting up the connection took over 20 seconds for the WRDS server. I infer that ADBC driver wants to know everything about the database it is connecting to before doing any work and it takes a lot longer for the WRDS database to tell its story than it does for my cheerful little home database.⁸

⁸A similar story is seen with connecting to ADBC databases with R, though it’s only nine seconds for WRDS to tell ADBC about itself. An issue with having an ADBC engine for the R version of `db2pq` is that the CRAN version of `adbcpostgresql` does not come with SSL support and I don’t want to be asking `db2pq` users to be compiling libraries against versions of `libpq`, etc.

The second thing I learnt is that the ADBC drivers are very conservative when it comes to mapping data types. For example, the `mthret` column from the `crsp.msf_v2` table that we saw above is stored in PostgreSQL as `NUMERIC(10, 6)`, which means its a precise representation of a number with 10 digits, 6 of which come after the decimal point. You might think that the ADBC driver would immediately think “hey, that’s the same thing as `pa.decimal128(10, 6)`” (here `pa` stands for `pyarrow`, the Python package for Arrow). But instead, the [driver documentation](#) for the ADBC driver for PostgreSQL states that “`NUMERIC` types are read as the string representation of the value, because the PostgreSQL `NUMERIC` type cannot be losslessly converted to the Arrow decimal types.”

In implementing the ADBC “engine”, I tweaked the implementation of functions such as `wrds_update_pq()` so that the one-time cost of setting up a database for connections with ADBC only happens once in a Python session. This is often fine even if that set-up takes more than 20 seconds because the Python session will run for minutes or even hours to download data.

I also gave the user the flexibility to specify how `NUMERIC` types are handled using the `numeric_mode` argument. With `numeric_mode="float64"` the values are cast by PostgreSQL to `DOUBLE PRECISION` before the ADBC driver sees them, so they end up as `Float64` types. If `numeric_mode="decimal"` the values are given to the ADBC driver as text, but these values are converted to nearest Arrow decimal types using the precision-scale metadata for the columns provided by PostgreSQL. Finally, `numeric_mode="text"` falls back to the default handling described by the Arrow documentation above.

3.3 Method 3: The `pg_parquet` PostgreSQL extension

The final contender is the `pg_parquet` PostgreSQL extension. The [Postgres.app site](#) offers a number of PostgreSQL extensions. These “are very easy to install – just download and double click the installer package! After restarting the server [also very easy with `Postgres.app`], you can enable them with the `CREATE EXTENSION` command.”

The `pg_parquet` PostgreSQL extension is one of several made available on the `Postgres.app` website. There it says: “Apache Parquet is gaining popularity as an interchange format for datasets. This extension allows you to read and write Parquet files using the `COPY` command. The extension was developed by CrunchyData and documentation is available at github.com/CrunchyData/pg_parquet.”

With this extension I can run SQL like this:

```
COPY crsp.msf_v2 TO '/tmp/msf_v2.parquet' (format 'parquet')
```

An obvious limitation of `pg_parquet` is that you need to have administrative privileges for the database you want to get data from. So clearly it’s not useful to a user of the WRDS database (unless WRDS were to install it, of course).

4 Benchmark analysis

Now that I have my three contenders, I run some benchmarks. I have copies of `comp.company`, `comp.funda`, and `crsp.msf_v2` in my local database and I use the code like `db_to_pq("msf_v2", "crsp", engine="duckdb")` to export the data to Parquet form for the first two contenders and `COPY` code like that above for `pg_parquet`. Each export includes all columns. Each bar represents a single export, with ADBC run twice for the small table to show first-run overhead. The benchmarks below were rerun on 8 October 2026 using the following versions.

Package or server	Version
db2pq	0.3.0
duckdb	1.5.6
adbc-driver-postgresql	1.12.0
adbc-driver-manager	1.12.0
pyarrow	25.0.1
pandas	3.0.6
polars	2.0.0
psycopg	3.3.6
matplotlib	3.11.2
seaborn	0.13.2
PostgreSQL	18.6 (Postgres.app)
pg_parquet	0.5.1

4.1 Small queries

The first export I look at is 50,000 rows from `comp.company`, which results in Parquet files of about 9 MiB. This small export allows me to demonstrate the cost of the first ADBC export within a Python instance. For the ADBC engine I run the query twice and you can see that the first time (`adbc_init` in Figure 1) is about five seconds longer than the second time (`adbc`). This difference includes initialization and connection setup; it is not a separate measurement of connection time.

In the original March 2026 tests, this setup took at least 20 seconds on the WRDS server. It seems that ADBC is doing a lot of “introspection” of the database each time a new database is established and this is a cost that is incurred for each run of Python (at least).




Figure 1: Export times for `comp.company`

4.2 Larger query, wide table

For the next query, I look at `comp.funda`, a table of financial statement information for (primarily) US companies. This table has over 900,000 rows and more than 900 variables, so it's a fairly wide table. The resulting Parquet files range from 405 to 454 MiB in size.

As can be seen in Figure 2, the DuckDB engine is fastest with this table at about 26 seconds, followed by ADBC at 55 seconds and `pg_parquet` at 189 seconds.




Figure 2: Export times for `comp.funda`

4.3 Larger query, long table

For the next query, I look at `crsp.msrf_v2`, a table of monthly data on the stock prices of US companies. This table has over 5,000,000 rows and 45 columns. It also has several columns that are stored in PostgreSQL as `NUMERIC` types (e.g., `mtthret`). The resulting Parquet files range from 219 to 231 MiB in size.

As can be seen in Figure 3, ADBC is more competitive with this table at about 15 seconds, compared with 12 seconds for the DuckDB engine and 53 seconds for `pg_parquet`.




Figure 3: Export times for `crsp.msf_v2`