

Data curation and the data science workflow

Ian D. Gow

6 July 2024

The goal of this note is to outline one part of extended version of model of the data science “whole game” proposed in [R for Data Science](#) (Wickham, Çetinkaya-Rundel, and Grolemund 2023). The original “whole game” comprises three steps. It starts with an *Import-and-tidy* process (this comprises *import* and *tidy*), then an *Understand* process (this involves iteration between *transform*, *visualize*, and *model* steps), followed by a *Communicate* process.¹

My extension of the “whole game”—depicted in Figure 1 below—adds a *persist* step to the *Import-and-tidy* process and re-labels this process as *Curate*. As a complement to the new *persist* step, I also add a *load* step to the *Understand* process. As will be seen this, *load* step will not generally be an elaborate one. The inclusion of a separate *load* step serves more to better delineate the distinction between the *Curate* process and the *Understand* process.

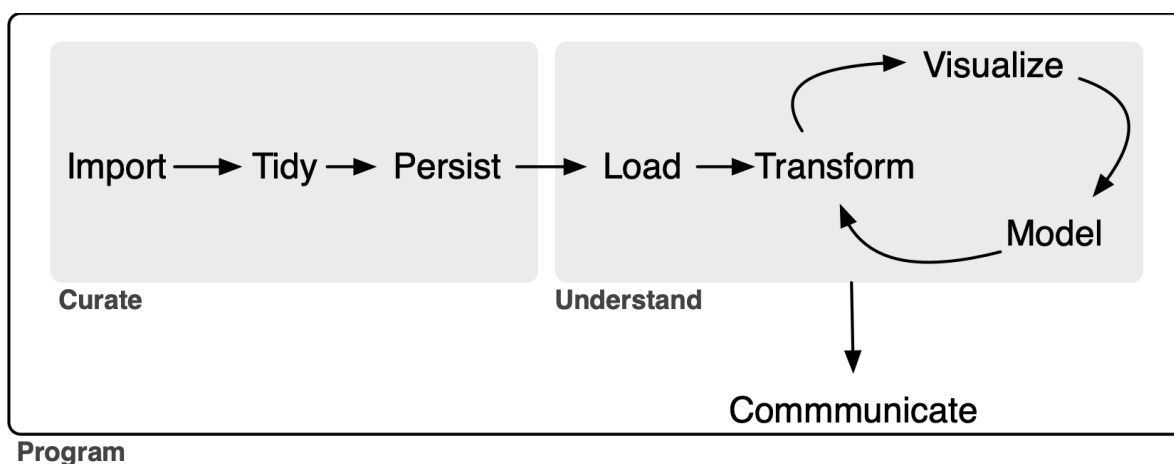


Figure 1: A representation of the data science workflow

¹The terms “process” and “step” are my own concoctions here and represent an attempt to group certain things together. I use capitalized verbs to describe what I am calling processes and lower-case verbs to denote steps. The original “whole game” model has just a single step after the *Understand* process and I upgrade that single step to a process.

In this note, I focus on the data curation (*Curate*) process. My rationale for separating *Curate* from *Understand* is that I believe it clarifies certain best practices in the curation of data. As will be seen, my conception of *Curate* will encompass some tasks that are included in the *transform* step (part of the *Understand* process) in [R for Data Science](#).

While I will argue that even the sole analyst who will perform all three processes can benefit from thinking about *Curate* separate from *Understand*, it is perhaps easiest to conceive of the *Curate* and *Understand* processes as involving different individuals or organizational units of the “whole game” of a data analysis workflow.²

This note was written using [Quarto](#) and compiled with [RStudio](#), an integrated development environment (IDE) for working with R. The source code for this note is available [here](#).

1 Possible data curation scenarios

In a *university* a “data lab” might be responsible for curating data sets for research and teaching, with researchers and students serving as the clients. Given the significance of data in a lot of research conducted in universities, my experience is that teams that might perform this role steadfastly refuse to do so. As a result, almost all data curation is performed by researchers themselves, often at great cost and with little expertise.³

In principle, *information vendors* are in the business of data curation. Some vendors do an excellent job of data curation and this makes it relatively easy for researchers to work with their data. For example, [CRSP](#) has long been the gold standard for researchers requiring data on US stock prices. For me, working with CRSP data is mostly a matter of getting it into my PostgreSQL database using [a script](#).⁴ Often, however, data vendors supply data in a way that requires significant additional data curation.

I imagine that larger quantitative *hedge funds* employ specialists to handle data curation tasks, with analysts and portfolio managers focusing on the tasks they are better suited to. No doubt there is iteration between data and IT specialists and those analysts and portfolio managers, as the output of the latter group needs to rigorously back-tested and put into production.

Research assistants might be hired by academic researchers to provide data curation services. In many cases, these research assistants will also do significant work from the *Understand* process, such as running regressions or other analyses. Even a *solo analyst* will be doing data curation, though the “client” will be the analyst herself.

I suspect that some of the principles I outline here will be useful in all of these scenarios.

²The authors of *R for Data Science* call the “whole game” a process, but I’ve already used that term to describe the next level down. So I choose *workflow* to denote the whole shebang here.

³Within business schools, my sense is that Chicago offers the best data curation support to faculty. While much of this occurs in a fairly decentralized fashion through the employment of [research professionals](#) who work closely with one or two faculty members, Chicago Booth also provides excellent research computing infrastructure, such as provisioning PostgreSQL databases for faculty use.

⁴More recently I have moved to also downloading CRSP data as parquet files using my [db2pq Python package](#).

2 The service-level agreement

It is perhaps helpful to think of dividing the data science workflow into processes with different teams being responsible for the different processes. From this perspective, the *Curate* team manufactures data that are delivered to the *Understand* team.⁵ While I won't discuss transfer pricing (i.e., how much the *Understand* team needs to pay the *Curate* team for the data), we might consider the analogy of a [service-level agreement](#) between the two teams.

One template for a service-level agreement would specify that data from a particular source will be delivered to the *Understand* team with the following conditions:

1. The data will be presented as a set of tables in a modern storage format.
2. The division into tables will adhere to a pragmatic version of good database principles.
3. The **primary key** of each table will be identified and validated.
4. Each variable (column) of each table will be of the correct type.
5. There will be no manual steps that cannot be reproduced.
6. A process for updating the curated data will be established.
7. The entire process will be documented in some way.
8. Some process for version control of data will be maintained.

2.1 Storage format

In principle, the storage format should fairly minor detail determined by the needs of the *Understand* team. For example, if the *Understand* team works in Stata or Excel, then perhaps they will want the data in some kind of Stata format or as Excel files. However, I think it can be appropriate to push back on notions that data will be delivered in form that involves downgrading the data or otherwise compromises the process in a way that may ultimately add to the cost and complexity of the task for the *Curate* team. For example, "please send the final data as an Excel file attachment as a reply email" might be a request to be resisted because the process of converting to Excel can entail the degradation of data (e.g., time stamps or encoding of text).⁶ Instead it may be better to choose a more robust storage format and supply a script for turning that into a preferred format.

One storage format that I have used in the past would deliver data as tables in a (PostgreSQL) database. The *Understand* team could be given access data from a particular source organized as a **schema** in a database. Accessing the data in this form is easy for any modern software package. One virtue of this approach is that the data might be curated using, say, Python even though the client will analyse it using, say, Stata.⁷

⁵By "manufacture" I merely mean to connote some notion of a production process, not some idea of [untoward processes for producing data](#).

⁶I discuss some of the issues with Excel as a storage format below.

⁷One project I worked on involved Python code analysing text and putting results in a PostgreSQL database and a couple of lines of code were sufficient for a co-author in a different city to load these data into Stata.

2.2 Good database principles

I included the word “pragmatic” because I think it’s not necessary in most cases to get particularly fussy about [database normalization](#). That said, it’s probably bad practice to succumb to requests for One Big Table that the *Understand* team might make. It is reasonable to impose some obligation to merge tables that are naturally different tables on the client *Understand* team.

2.3 Primary keys

The *Curate* team should communicate the primary key of each table to the *Understand* team.⁸ A primary key of a table will be a set of variables that can be used to uniquely identify a row in that table. In general a primary key will have no missing values. Part of data curation will be confirming that a proposed primary key is in fact a valid primary key.

2.4 Data types

Each variable of each table should be of the correct type. For example, dates should be of type DATE, variables that only take integer values should be of INTEGER type.⁹ Date-times should generally be given with `TIMESTAMP WITH TIME ZONE` type. [Logical columns](#) should be supplied with type BOOLEAN.

Note that there is an interaction between this element of the service-level agreement and the storage format. If the data are supplied in a PostgreSQL database or as parquet files, then it is quite feasible to prescribe the data types of each variable. But if the storage format is Excel files (not recommended!) or CSV files, then it is difficult for the data curator to control how each variable is understood by the *Understand* team.¹⁰

In some cases, it may seem unduly prescriptive to specify the types in a particular way. For example, a logical variable can easily be represented as INTEGER type (0 for FALSE, 1 for TRUE). Even in such cases, I think there is merit in choosing the most logical type (no pun intended) because of the additional information it conveys about the data. For example, a logical type should be checked to ensure that it only takes two values (TRUE or FALSE) plus perhaps NULL and that this checking has occurred is conveyed by the encoding of that variable as BOOLEAN.

⁸Sometimes there will be more than one primary key for a table.

⁹Here I used PostgreSQL data types, but the equivalent types in other formats should be fairly clear.

¹⁰SAS and Stata are somewhat loose with their “data types”. In effect SAS has just two data types—fixed-width character and floating-point numeric—and the other types are just formatting overlays over these. These types can be easily upset depending on how the data are used.

2.5 No manual steps

When data vendors are providing well-curated data sets, much about the curation process will be obscure to the user. This makes some sense, as the data curation process has elements of trade secrets. But often data will be supplied by vendors in an imperfect state and significant data curation will be performed by the *Curate* team working for or within the same organization as the *Understand* team.

Focusing on the case where the data curation process transforms an existing data set—say, one purchased from an outside vendor—into a curated data set in sense used here, there are a few ground rules regarding manual steps.

First, *the original data files should not be modified in any way*. If data are supplied as CSV files, then merely opening them in Excel and saving them can mutilate the original data.¹¹ I have encountered people whose idea of data curation extended to opening the original files, saving them as Excel files, and then proceeding to manually edit those files. This approach leads to a completely unreproducible set of data files, which is problematic not only in a world in which reproducibility is starting to be expected, but also when a new version of the data will be supplied by the vendor in the future.

Second, any manual steps should be extensively documented and applied in a transparent automated fashion. For example, if a data set on financial statement items of US firms contains errors that can be corrected by reviewing original SEC filings, then any corrections should be clearly documented in separate files with links to the original filings and explanations. And the corrections should be implemented through code, not manual steps. For example, there should be code that imports the original data and the corrections and applies the latter to the former to create the final data set.

2.6 Documentation

The process of curating the data should be documented sufficiently well that someone else could perform the curation steps should the need arise. Often that need will arise when the vendor provides an updated data set. Perhaps the best way to understand what I have in mind here is through a case study and I provide one in Section 3.

2.7 Update process

Part of the rationale for having a well-documented process with no manual steps is that it greatly facilitates updating the data when the data vendor or other data source provides updated data. In some cases, updating the data will entail little more than downloading the new raw data and running a pre-existing script on those data. In other cases, the data may change in significant

¹¹An example where this could happen is provided in Tip 1.

ways, such as addition of new variables, renaming of existing ones, or reorganization of data into different tables.

As future changes may be difficult to predict, the analyst might be able to do little more than describe the anticipated update process if no major changes occur. If major changes do subsequently occur, it likely makes sense for the analyst handling the update to extensively document the changes needed to process the new data, especially if earlier versions of the data remain relevant (e.g., they have been used in published research).

2.8 Data version control

Welch (2019) argues that, to ensure that results can be reproduced, “the author should keep a private copy of the full data set with which the results were obtained.” This imposes a significant cost on the *Understand* team to maintain archives of data sets that may run to several gigabytes or more and it would seem much more efficient for these obligations to reside with the parties with the relevant expertise.

Unfortunately, even when data vendors provide curated data sets, they generally provide little in the way of version control. For example, there is no evidence that Wharton Research Data Services (WRDS), perhaps the largest data vendor in academic business research, provides any version control for its datasets, even though it should have much greater expertise for doing this than the users of its services.

Nonetheless, some notion of version control of data probably has a place in data curation, even if this is little more than archiving of various versions of data supplied to research teams.

3 A data curation case study: SIRCA ASX EOD data

In this section, I use the packages listed below, plus the `duckdb` package.¹² This note was written using [Quarto](#) and compiled with [RStudio](#), an integrated development environment (IDE) for working with R. The source code for this note is available [here](#).

```
library(tidyverse)
library(DBI)
library(arrow)
library(farr)
```

The tables included with the SIRCA ASX EOD price collection are listed in Table 2. Each of these tables is supplied by SIRCA in the form of a compressed comma-separated values (CSV) file. For example, `si_au_ref_names` is supplied as `si_au_ref_names.csv.gz`.

¹²Execute `install.packages(c("tidyverse", "DBI", "duckdb", "arrow", "farr"))` within R to install all the packages you need to run the code in this note.

The first step of our process will be to obtain these four CSV files and save them in a subdirectory named `sirca` on your computer. You should specify the location of that subdirectory by editing the following command, replacing with `"~/Library/CloudStorage/Dropbox/raw_data/"` with, say, `"C:\Data\CSV Files"`, if that is where you have created this `sirca` directory on your computer.

```
Sys.setenv(RAW_DATA_DIR = "~/Library/CloudStorage/Dropbox/raw_data")
```

Thus, the CSV files should be stored in the location that we will assign to the variable `csv_dir`.

```
csv_dir <- file.path(Sys.getenv("RAW_DATA_DIR"), "sirca")
```

From Table 1, we can see that we have three files of fairly modest size and one large file (`si_au_prc_daily.csv.gz`). Note that the largest file will be about 10 times larger when decompressed. Because larger files present their own issues, we will start with `si_au_ref_names`, which presents some complexity while being fairly easy to work with.

Table 1: Data on supplied CSV files from SIRCA

| file_name | size | mtime |
|---|-----------|---------------------|
| <code>si_au_prc_daily.csv.gz</code> | 365.44 MB | 2024-06-28 21:19:04 |
| <code>si_au_ref_names.csv.gz</code> | 591.81 kB | 2024-06-28 21:22:06 |
| <code>si_au_ref_trddays.csv.gz</code> | 59.65 kB | 2024-06-29 06:07:05 |
| <code>si_au_retn_mkt.csv.gz</code> | 352.44 kB | 2024-06-29 06:07:09 |
| <code>SIRCA EOD Data Dictionary.xlsx</code> | 174.3 kB | 2024-06-29 06:06:50 |

Table 2: SIRCA ASX EOD price collection

| Table | Description |
|--------------------------------|---|
| <code>si_au_ref_names</code> | Name histories and change dates for companies listed since 1 January 2000 |
| <code>si_au_prc_daily</code> | Complete daily price, volume and value histories, with issued share numbers |
| <code>si_au_retn_mkt</code> | Daily value- and equal-weighted whole-market returns |
| <code>si_au_ref_trddays</code> | Record of ASX trading dates since 1 January 2000 |

Here we choose parquet files as our target storage format. We will store our data in a `sirca` subdirectory in a different location from `RAW_DATA_DIR` specified above. You should specify

the location DATA_DIR by editing the line of code below, much as you specified RAW_DATA_DIR above.¹³

```
Sys.setenv(DATA_DIR = "~/Library/CloudStorage/Dropbox/pq_data/")
```

3.1 Importing si_au_ref_names

As discussed above, we start with si_au_ref_names. We first specify the name of the CSV file si_au_ref_names_csv, then quickly move on to reading the data using the read_csv() function. The displayed output from invoking read_csv() provides a good starting point for the next steps.

3.1.1 Setting data types

As can be seen, si_au_ref_names contains 20 columns that read_csv() parses as character columns and 9 columns that read_csv() parses as numeric columns.

```
si_au_ref_names_csv <- file.path(csv_dir, "si_au_ref_names.csv.gz")
si_au_ref_names <- read_csv(si_au_ref_names_csv)
```

```
Rows: 11679 Columns: 29
-- Column specification -----
Delimiter: ","
chr (20): Gcode, CompanyTicker, SecurityTicker, SecurityType, Abreviate...
dbl (9): SeniorSecurity, ListDate_YMD, DelistDate_YMD, ListDate_DaysSi...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

The next step we take is to inspect the columns to determine whether refinement of types makes sense. In practice, we can infer appropriate types by looking at the data.

We start with three of the numeric columns. The first three appear to be integers, either based on casual inspection of the values displayed or inferences from the variable names (e.g., “days since” seems likely to be an integer).

¹³Using environment variables to specify RAW_DATA_DIR and DATA_DIR may not have an obvious payoff in the context of this note. The benefit comes more from follow-on work using the data and also from applying the approach to managing raw data more broadly.


```
si_au_ref_names |>
  select_if(is.numeric) |>
  select(1, 4:5)
```

```
# A tibble: 11,679 x 3
  SeniorSecurity ListDate_DaysSince DelistDate_DaysSince
      <dbl>           <dbl>           <dbl>
1         1         44074             NA
2         1         43355             NA
3         1         42604             NA
4         1         44699             NA
5         1         45197             NA
# i 11,674 more rows
```

We can check that converting these variables to integers using `as.integer()` does not change any of their values.

```
si_au_ref_names |>
  select_if(is.numeric) |>
  select(1, 4:5) |>
  summarize(across(everything(), \(x) all(x == as.integer(x), na.rm = TRUE))))
```

```
# A tibble: 1 x 3
  SeniorSecurity ListDate_DaysSince DelistDate_DaysSince
      <lgl>           <lgl>           <lgl>
1 TRUE          TRUE          TRUE
```

We can do the same for four of the remaining numeric columns.

```
si_au_ref_names |>
  select_if(is.numeric) |>
  select(6:9) |>
  summarize(across(everything(), \(x) all(x == as.integer(x), na.rm = TRUE))))
```

```
# A tibble: 1 x 4
  RecordCount GICSIndustry SIRCAIndustryClassCode SIRCASectorCode
      <lgl>      <lgl>      <lgl>           <lgl>
1 TRUE      TRUE      TRUE          TRUE
```

The remaining numeric variables appear to be dates in ymd form read by `read_csv()` as numeric variables.

```
si_au_ref_names |>
  select_if(is.numeric) |>
  select(2:3)
```

```
# A tibble: 11,679 x 2
  ListDate_YMD DelistDate_YMD
      <dbl>         <dbl>
1      20200831             NA
2      20180912             NA
3      20160822             NA
4      20220518             NA
5      20230928             NA
# i 11,674 more rows
```

We can convert these columns to dates with the `ymd()` function. In the following code snippet, we convert the numeric variables to the types we determined to be appropriate through the analysis above. Here this code just tests that nothing untoward happens; we will actually implement these type conversions in code below.

```
si_au_ref_names |>
  select_if(is.numeric) |>
  mutate(across(c(SeniorSecurity, ListDate_DaysSince, DelistDate_DaysSince,
                  RecordCount, GICSIndustry, SIRCAIndustryClassCode,
                  SIRCASectorCode), as.integer),
         across(ends_with("_YMD"), ymd))
```

```
# A tibble: 11,679 x 9
  SeniorSecurity ListDate_YMD DelistDate_YMD ListDate_DaysSince
      <int> <date>         <date>         <int>
1          1 2020-08-31      NA             44074
2          1 2018-09-12      NA             43355
3          1 2016-08-22      NA             42604
4          1 2022-05-18      NA             44699
5          1 2023-09-28      NA             45197
# i 11,674 more rows
# i 5 more variables: DelistDate_DaysSince <int>, RecordCount <int>,
#   GICSIndustry <int>, SIRCAIndustryClassCode <int>,
#   SIRCASectorCode <int>
```

We can now move onto the 20 columns read as character vectors. The first five character vectors seem correctly identified as such.

```
si_au_ref_names |>
  select_if(is.character) |>
  select(1:5)
```

```
# A tibble: 11,679 x 5
  Gcode CompanyTicker SecurityTicker SecurityType AbbreviatedSecurityDescr~1
  <chr> <chr>          <chr>          <chr>          <chr>
1 2181 218            218            01            nsx - ordinary
2 14d1 14D            14D            01            ordinary
3 1ad1 1AD            1AD            01            ordinary
4 1ae1 1AE            1AE            01            ordinary
5 1gov1 1GO          1GOV           07            etf units
# i 11,674 more rows
# i abbreviated name: 1: AbbreviatedSecurityDescription
```

The same is true for character vectors 8 and 9 ...

```
si_au_ref_names |>
  select_if(is.character) |>
  select(8:9)
```

```
# A tibble: 11,679 x 2
  FullCompanyName AbbrevCompanyName
  <chr>           <chr>
1 ROFINA GROUP LIMITED ROFINA GROUP LIMITED
2 1414 DEGREES LIMITED 1414 DEGREES LIMITED
3 ADALTA LIMITED      ADALTA LIMITED
4 AURORA ENERGY METALS LIMITED AURORAENERGYMETALS
5 VANECK 1-5 YEAR AUSTRALIAN GOVERNMENT BOND ETF VANECK 1-5 YR GOV
# i 11,674 more rows
```

... and character vectors 14 through 15 ...

```
si_au_ref_names |>
  select_if(is.character) |>
  select(14:15) |>
  filter(if_all(everything(), \(x) !is.na(x)))
```

```
# A tibble: 22 x 2
  CompanyDelistReasonComment AlteredLink
  <chr>                      <chr>
1 converts to a trust by a one-for-one in specie issue in trus~ [aqf] is o~
2 pursuant to scheme of arrangement with arrow pharmaceuticals~ {awp2}[awp~
3 <18/02/2000>. Demerger. {bor1} (bor) boral limited split int~ {bor1} (bo~
4 redomiciled to New Zealand after one for one share exchange ~ redomicile~
5 <05/04/2002> discount card shares consolidated with ordinary~ first 500 ~
# i 17 more rows
```

... and character vectors 16 through 20.

```
si_au_ref_names |>
  select_if(is.character) |>
  select(16:20) |>
  filter(if_all(everything(), \(x) !is.na(x)))
```

```
# A tibble: 26 x 5
  MS_CompanyID MS_SecurityID MS_CompanyID2 MS_SecurityID2 MA_Identifier
  <chr>         <chr>         <chr>         <chr>         <chr>
1 OC00000ALG   OP0000A1JX     OC00008X4T    OP0001D0H2     CDE
2 OC00000ALG   OP0000A1JX     OC00008X4T    OP0001D0H2     CXC
3 OC00000ALG   OP0000A1JX     OC00008X4T    OP0001D0H2     CXC
4 OC00000ALG   OP0000A1JX     OC00008X4T    OP0001D0H2     CXC
5 OC00000LFF   OP00006XWH     OC00009CVF    OP0000WDB0     SW1
# i 21 more rows
```

This leaves character columns 12 and 13. Focusing on the cases where neither is NA, we see that these columns appear to be lists of codes separated by semi-colons (;).

```
si_au_ref_names |>
  select_if(is.character) |>
  select(12:13) |>
  filter(if_all(everything(), \(x) !is.na(x)))
```

```
# A tibble: 477 x 2
  CompanyDelistReasonCode CompanyRelatedGCode
  <chr>                  <chr>
1 A                      tail
2 A                      wgr1
3 N                      aln2; agl1; agk1
```

```

4 R;A          bep1; bnb1; pif1; bbp1; bbw1
5 A           btx1
# i 472 more rows

```

```

si_au_ref_names |>
  mutate(CompanyRelatedGCodes = CompanyRelatedGCode) |>
  filter(CompanyRelatedGCode != "") |>
  mutate(related_gcode = str_split(CompanyRelatedGCode, "[;\s]+")) |>
  select(Gcode, CompanyRelatedGCode, related_gcode)

```

```

# A tibble: 480 x 3
  Gcode CompanyRelatedGCode      related_gcode
  <chr> <chr>                <list>
1 a1c1  tai1                  <chr [1]>
2 aag2  wgr1                  <chr [1]>
3 aan2  aln2; agl1; agk1      <chr [3]>
4 aan2  bep1; bnb1; pif1; bbp1; bbw1 <chr [5]>
5 aby1  btx1                  <chr [1]>
# i 475 more rows

```

```

si_au_ref_names |>
  filter(CompanyRelatedGCode != "") |>
  mutate(related_gcode = str_split(CompanyRelatedGCode, "[;\s]+")) |>
  unnest(related_gcode) |>
  select(Gcode, CompanyRelatedGCode, related_gcode)

```

```

# A tibble: 514 x 3
  Gcode CompanyRelatedGCode related_gcode
  <chr> <chr>                <chr>
1 a1c1  tai1                  tai1
2 aag2  wgr1                  wgr1
3 aan2  aln2; agl1; agk1      aln2
4 aan2  aln2; agl1; agk1      agl1
5 aan2  aln2; agl1; agk1      agk1
# i 509 more rows

```

Tip 1: Friends don't let friends use Excel

From casual observation, it appears that valid Gcode values contain only lower case characters ([a-z] in regular expressions) or numbers ([0-9] in regular expressions). Are there any CompanyRelatedGCode values that contain other characters? It turns out that there

are.

```
si_au_ref_names |>
  mutate(related_gcode = str_split(CompanyRelatedGCode, "[;\\s]+")) |>
  unnest(related_gcode) |>
  filter(str_detect(related_gcode, "[^a-z0-9]")) |>
  select(Gcode, related_gcode, CompanyDelistReasonComment)
```

```
# A tibble: 2 x 3
  Gcode related_gcode CompanyDelistReasonComment
  <chr> <chr>          <chr>
1 ahx1 May-01          mayne nickless ltd
2 fhf1 May-01          mayne nickless limited
```

What's happened here? May-01 looks more like a date than a Gcode. This has all the hallmarks of someone having imported data into Microsoft Excel as part of their process. Microsoft Excel has a well-known tendency to mangle values that it aggressively interprets as dates. It seems likely that the Gcode for Mayne Nickless was may1 and Excel read this as May-01 (a date).¹⁴ Is it true that Gcode values contain only lower-case characters and numbers?

```
weird_gcodes <-
  si_au_ref_names |>
  filter(str_detect(Gcode, "[^a-z0-9]")) |>
  distinct(Gcode)
```

It seems not; some Gcodes have underscores (_):

```
str_flatten(pull(weird_gcodes), ", ")
```

```
[1] "92e_1, apr_1, aug_1, aug_3, mar_2, may_1, nov_1, oct_1"
```

To see why underscores are used, we can remove the underscore and save the Gcodes in a CSV file.¹⁵

```
weird_gcodes |>
  mutate(Gcode = str_remove(Gcode, "_")) |>
  write_csv("weird_gcodes.csv")
```

Try opening weird_gcodes.csv in Excel. What do you see? (It may help to open weird_gcodes.csv in a text editor to see the original values.) To be frank, I struggle to see any reason why Excel should have any part in the data science workflow.¹⁶

We can examine `CompanyDelistReasonCode` in much the same way we did `CompanyRelatedGCode`. For reasons of brevity, I spare you the coding details and focus on the processed data, information about which is shown in Table 3.

One problem is evident from Table 3 and that is the presence of what appears to be junk in the `CompanyDelistReasonCode` field (e.g., 18 or R-apx). Another problem is evident only after looking at the documentation for `si_au_ref_names` and that is that even when the codes appear well-formed (e.g., N or C), we have no information about what these codes mean.

Table 3: Delisting reason codes on `si_au_ref_names`

| delist_code | n | delist_code | n |
|-------------|------|-------------|----|
| N | 3855 | G | 15 |
| C | 3482 | X | 13 |
| R | 976 | I | 10 |
| A | 720 | T | 6 |
| S | 552 | Z | 6 |
| M | 371 | 18 | 5 |
| F | 279 | 2 | 5 |
| E | 273 | 9 | 5 |
| Y | 186 | D | 5 |
| W | 48 | P | 5 |
| L | 31 | B | 4 |
| H | 26 | p | 2 |
| U-x | 25 | 0 | 1 |
| O | 18 | R-apx | 1 |

Given the issues apparent in both `CompanyRelatedGCode` and `CompanyDelistReasonCode`, I have elected to collect those, but keep them as simple character columns.

For those keeping track, we have four character columns left. It turns out that the name for each of these ends with `Date`. In the following, I focus on the observations with non-NA values in all of these columns.

```
si_au_ref_names |>
  select_if(is.character) |>
  select(ends_with("Date")) |>
  filter(if_all(everything(), \(x) !is.na(x)))
```

¹⁶Whoever did this also had their computer set to format dates in the US-style `Mmm-dd` format, rather than the `dd-Mmm` style I see on my computer.

¹⁶You can download this CSV file [here](#).

¹⁶See Broman and Woo (2018) for further discussion of some of the issues with using Excel for data science.

```
# A tibble: 5,053 x 4
  ListDate   DelistDate EarliestListDate LatestDelistDate
  <chr>      <chr>      <chr>              <chr>
1 14/06/2023 29/08/2023 2/03/2021          29/08/2023
2 4/03/2019 10/01/2022 4/03/2019          10/01/2022
3 9/10/2015 1/03/2017 9/10/2015          1/03/2017
4 7/11/1997 9/11/2000 7/11/1997          9/11/2000
5 24/11/2020 24/04/2023 24/11/2020         24/04/2023
# i 5,048 more rows
```

From the above, it seems clear that we have dates in dmy form. It turns out that a couple of observations have the value "0/01/1900", which is not a valid date and I convert these to missing values using the code below.

```
si_au_ref_names |>
  select(ends_with("Date")) |>
  mutate(across(ends_with("Date"),
    \ (x) dmy(if_else(x == "0/01/1900", NA, x))))
```

```
# A tibble: 11,679 x 4
  ListDate   DelistDate EarliestListDate LatestDelistDate
  <date>      <date>      <date>              <date>
1 2020-08-31 NA          2020-08-31          NA
2 2018-09-12 NA          2018-09-12          NA
3 2016-08-22 NA          2016-08-22          NA
4 2022-05-18 NA          2022-05-18          NA
5 2023-09-28 NA          2023-09-28          NA
# i 11,674 more rows
```

At this point, we have two versions of the variables related to listing dates (ListDate_YMD and ListDate) and to delisting dates (DelistDate_YMD and DelistDate) and perhaps it makes sense to keep just one of each. If the values in each of the pair is the same as the other, then there's no reason to keep both.

Looking at ListDate_YMD and ListDate are always equal and we could drop either one and keep the other.

```
si_au_ref_names |>
  select(matches("ListDate")) |>
  mutate(across(ends_with("Date"),
    \ (x) dmy(if_else(x == "0/01/1900", NA, x))),
    across(ends_with("_YMD"), ymd)) |>
  filter(ListDate_YMD != ListDate)
```



```
# A tibble: 0 x 8
# i 8 variables: ListDate_YMD <date>, DelistDate_YMD <date>,
#   ListDate_DaysSince <dbl>, DelistDate_DaysSince <dbl>, ListDate <date>,
#   DelistDate <date>, EarliestListDate <date>, LatestDelistDate <date>
```

But there is one instance where DeistDate_YMD and DelistDate differ.

```
si_au_ref_names |>
  select(matches("^DelistDate")) |>
  mutate(across(ends_with("Date"),
    \(x) dmy(if_else(x == "0/01/1900", NA, x))),
    across(ends_with("_YMD"), ymd)) |>
  filter(DelistDate_YMD != DelistDate)
```

```
# A tibble: 1 x 3
  DelistDate_YMD DelistDate_DaysSince DelistDate
  <date>                <dbl> <date>
1 2023-06-05          45051 2023-05-05
```

Which one to choose? One approach would be to look to external sources to verify which date is correct. But for present purposes we will choose the one that keeps our data internally consistent. Specifically, we should choose whichever of DelistDate_YMD and DelistDate that is consistent with DelistDate_DaysSince.

Looking for other rows where DelistDate_DaysSince == 45051, we see that that value is elsewhere consistent with the value in DelistDate, so here I choose to drop the _YMD variables.

```
si_au_ref_names |>
  select(Gcode, starts_with("DelistDate")) |>
  filter(DelistDate_DaysSince == 45051)
```

```
# A tibble: 2 x 4
  Gcode DelistDate_YMD DelistDate_DaysSince DelistDate
  <chr>      <dbl>                <dbl> <chr>
1 iesg1    20230505          45051 5/05/2023
2 sur1     20230605          45051 5/05/2023
```

Putting all the pieces above we have the following:

```

si_au_ref_names <-
  read_csv(si_au_ref_names_csv, show_col_types = FALSE) |>
  mutate(across(c(SeniorSecurity, ListDate_DaysSince, DelistDate_DaysSince,
                  RecordCount, GICSIndustry, SIRCAIndustryClassCode,
                  SIRCASectorCode), as.integer),
         across(ends_with("Date"),
                \(x) dmy(if_else(x == "0/01/1900", NA, x)))) |>
  select(-ends_with("_YMD"))

```

3.1.2 Identifying the primary key

Before considering possible primary keys, we first determine if there are any duplicate rows. When there are duplicate rows, no possible combination of columns will work as a primary key.

The following function returns any rows that are duplicated in a data set.

```

get_dupes <- function(df, count_var = "count") {
  df |>
    count(pick(everything()), name = count_var) |>
    filter(.data[[count_var]] > 1)
}

```

Applying this function to `si_au_ref_names`, we see that we have one row that appears twice in the data set.

```

si_au_ref_names |>
  get_dupes() |>
  select(Gcode, SecurityTicker, ListDate, count)

```

```

# A tibble: 1 x 4
  Gcode SecurityTicker ListDate    count
  <chr> <chr>          <date>    <int>
1 oct_1 OCT            2022-09-16      2

```

To address this, we will simply use the `distinct()` function.

Moving on to consider potential primary keys, we see immediately that `(Gcode, SecurityTicker)` is not a valid primary key. As seen in the output below, a given `(Gcode, SecurityTicker)` combination can appear as many as seven times in the data.

```
si_au_ref_names |>
  distinct() |>
  count(Gcode, SecurityTicker, name = "num_rows") |>
  count(num_rows)
```

```
# A tibble: 7 x 2
  num_rows      n
  <int> <int>
1         1 6177
2         2 1678
3         3  450
4         4  127
5         5   44
6         6   10
7         7    1
```

Looking across the columns, we see that (Gcode, SecurityTicker, ListDate) *almost* works, as we have just one case where (Gcode, SecurityTicker, ListDate) fails to identify a single row. In this particular case, it seems that we have differences only in GICSIndustry and SIRCAIndustryClassCode. In one row, these variables are missing; in the other there are values.

```
si_au_ref_names |>
  distinct() |>
  group_by(Gcode, SecurityTicker, ListDate) |>
  filter(n() > 1) |>
  ungroup() |>
  arrange(Gcode, SecurityTicker, ListDate) |>
  select(Gcode, SecurityTicker, ListDate, GICSIndustry, SIRCAIndustryClassCode)
```

```
# A tibble: 2 x 5
  Gcode SecurityTicker ListDate   GICSIndustry SIRCAIndustryClassCode
  <chr> <chr>          <date>         <int>          <int>
1 rgwb1 RGWB          2007-05-01         NA             NA
2 rgwb1 RGWB          2007-05-01    999999999      26
```

If we take the row with non-NA values for GICSIndustry and SIRCAIndustryClassCode to be the correct one, then we should delete the other row.

```
si_au_ref_names |>
  filter(Gcode == "rgwb1") |>
  select(Gcode, GICSIndustry)
```

```
# A tibble: 2 x 2
  Gcode GICSIndustry
  <chr>      <int>
1 rgwb1         NA
2 rgwb1      99999999
```

It turns out that these are the only two rows where `Gcode == "rgwb1"`, so if we eliminate the row with NA value in `GICSIndustry` we should have it that `(Gcode, SecurityTicker, ListDate)` uniquely identifies each row.

```
si_au_ref_names |>
  distinct() |>
  filter(!(Gcode == "rgwb1" & is.na(GICSIndustry))) |>
  count(Gcode, SecurityTicker, ListDate, name = "num_rows") |>
  count(num_rows)
```

```
# A tibble: 1 x 2
  num_rows      n
  <int> <int>
1       1 11677
```

To confirm that `(Gcode, SecurityTicker, ListDate)` is a valid primary key for our filtered `si_au_ref_names`, we also need to check that there are no NA values in any of these fields, which the following code confirms.

```
si_au_ref_names |>
  distinct() |>
  filter(!(Gcode == "rgwb1" & is.na(GICSIndustry))) |>
  summarize(across(c(Gcode, SecurityTicker, ListDate),
    \ (x) all(!is.na(x))
```

```
# A tibble: 1 x 3
  Gcode SecurityTicker ListDate
  <lgl> <lgl>          <lgl>
1 TRUE  TRUE          TRUE
```

3.1.3 Writing the parquet file

So, we can put the reading of raw data, the conversion of data types, and the filters needed to have a valid primary key together. But we have one final adjustment to make and that is to convert all variable names to lower case, as we will see later that the variable names embedded in `si_au_prc_daily.csv.gz` are all lower case (e.g., `gcode`), so we probably make our lives easier by converting our variables here to lower case (e.g., so we can join on `gcode` without worrying about slight differences in variable names).

With that final adjustment, we can then write to a parquet file, as we do here. We will use the environment variable `DATA_DIR` that you set above to specify the location.

```
pq_dir <- file.path(Sys.getenv("DATA_DIR"), "sirca")
if (!dir.exists(pq_dir)) dir.create(pq_dir)

si_au_ref_names <-
  read_csv(si_au_ref_names_csv, show_col_types = FALSE) |>
  mutate(across(c(SeniorSecurity, ListDate_DaysSince, DelistDate_DaysSince,
                  RecordCount, GICSIndustry, SIRCAIndustryClassCode,
                  SIRCASectorCode), as.integer),
         across(ends_with("Date"),
                \(x) dmy(if_else(x == "0/01/1900", NA, x)))) |>
  select(-ends_with("_YMD")) |>
  distinct() |>
  filter(!(Gcode == "rgwb1" & is.na(GICSIndustry))) |>
  rename_with(str_to_lower) |>
  write_parquet(sink = file.path(pq_dir, "si_au_ref_names.parquet")) |>
  system_time()
```

```
user  system elapsed
0.083  0.012  0.083
```

3.2 Importing `si_au_ref_trddays`

A similar process to that used for `si_au_ref_names` can be applied to `si_au_ref_trddays`. However, `si_au_ref_trddays` is a much simpler file and we conclude that the types of the five columns can be specified using `col_types = "ciDii"`, where `c` means character, `i` means integer, and `D` means date.¹⁷

¹⁷See the help for `read_csv()` to learn more.

```
si_au_ref_trddays_csv <- file.path(csv_dir, "si_au_ref_trddays.csv.gz")

si_au_ref_trddays <-
  read_csv(si_au_ref_trddays_csv, col_types = "ciDii") |>
  mutate(dateymd = ymd(dateymd))
```

We can easily confirm that date is a valid primary key:

```
si_au_ref_trddays |>
  count(date, name = "num_rows") |>
  count(num_rows)
```

```
# A tibble: 1 x 2
  num_rows      n
    <int> <int>
1         1 6072
```

```
si_au_ref_trddays |>
  summarize(across(date, \(x) all(!is.na(x))))
```

```
# A tibble: 1 x 1
  date
  <lgl>
1 TRUE
```

We can also confirm that we don't need dateymd, as it contains the same information as date.

```
si_au_ref_trddays |>
  filter(dateymd != date) |>
  count() |>
  pull()
```

```
[1] 0
```

We can specify - in col_types to omit dateymd when we read the data. Since date will be our primary key, we put that column first using the relocate() function.

```
si_au_ref_trddays <-
  read_csv(si_au_ref_trddays_csv, col_types = "-iDii") |>
  relocate(date)
```

We also confirm that `dayssince` simply represents the number of dates since 1899-12-30.

```
si_au_ref_trddays |>
  mutate(some_date = date - dayssince) |>
  count(some_date)
```

```
# A tibble: 1 x 2
  some_date      n
  <date>      <int>
1 1899-12-30  6072
```

We can also confirm that `weekday` represents the day of the week in the US system that starts the week on Sunday.¹⁸

```
si_au_ref_trddays |>
  mutate(
    weekday_calc = wday(date),
    wday = wday(date, label = TRUE)
  ) |>
  count(weekday, weekday_calc, wday)
```

```
# A tibble: 5 x 4
  weekday weekday_calc wday      n
  <int>      <dbl> <ord> <int>
1      2          2 Mon     1165
2      3          3 Tue     1225
3      4          4 Wed     1234
4      5          5 Thu     1235
5      6          6 Fri     1213
```

```
si_au_ref_trddays <-
  read_csv(si_au_ref_trddays_csv,
    col_types = "-iDii") |>
  relocate(date) |>
  write_parquet(sink = file.path(pq_dir, "si_au_ref_trddays.parquet")) |>
  system_time()
```

```
user  system elapsed
0.017  0.003   0.018
```

¹⁸The ISO 8601 convention is more consistent with the idea that Sunday is at the *end* of the week—hence “weekend”—and starts the week on Monday. But these distinctions are not important here.

3.3 Importing si_au_retn_mkt

We omit the details, but we can confirm that much of what we saw with `si_au_ref_trddays` applies to `si_au_retn_mkt`:

- Date is a valid primary key
- DateYMD is redundant
- DaysSince represents the number of days since 1899-12-30

Again we convert all column names to lower case so that date is a common field across `si_au_ref_trddays`, `si_au_retn_mkt`, and `si_au_prc_daily`.

```
si_au_retn_mkt_csv <- file.path(csv_dir, "si_au_retn_mkt.csv.gz")

si_au_retn_mkt <-
  read_csv(si_au_retn_mkt_csv,
           col_types = "-iDdddddd",
           locale = locale(date_format = "%d/%m/%Y"),
           name_repair = str_to_lower) |>
  relocate(date) |>
  write_parquet(sink = file.path(pq_dir, "si_au_retn_mkt.parquet")) |>
  system_time()
```

```
user  system elapsed
0.015  0.005   0.014
```

3.4 Importing si_au_prc_daily

By this point, we should be getting the hang of the workflow. We now move on to the largest file in the set, `si_au_prc_daily.csv.gz`. We start by identifying the CSV source and the parquet destination.

```
si_au_prc_daily_csv <- file.path(csv_dir, "si_au_prc_daily.csv.gz")
si_au_prc_daily_pq <- file.path(pq_dir, "si_au_prc_daily.parquet")
```

Using a process similar to that above, we identify those columns needing special handling in the import process. Note that we specify `guess_max = 1e6` because the default value for `guess_max` reads too few rows to infer the types of some variables that are most NA.


```

si_au_prc_daily <-
  read_csv(si_au_prc_daily_csv,
           guess_max = 1e6,
           show_col_types = FALSE) |>
  mutate(dateymd = ymd(dateymd),
         date = dmy(date),
         weekday = as.integer(weekday),
         monthend = as.logical(monthend),
         seniorsecurity = as.integer(seniorsecurity)) |>
  system_time()

```

```

user  system elapsed
94.494   7.861  58.675

```

Again we need to choose between date and dateymd, which are almost always equal.

```

si_au_prc_daily |>
  filter(date != dateymd) |>
  select(gcode, securityticker, date, dateymd, dayssince)

```

```

# A tibble: 1 x 5
  gcode securityticker date      dateymd    dayssince
<chr> <chr>          <date>    <date>      <dbl>
1 ind1  PMXDA          2011-05-09 2011-05-19    40682

```

Again dateymd seems to be the one of the two that is consistent with dayssince.

```

si_au_prc_daily |>
  filter(dayssince == 40682) |>
  count(date)

```

```

# A tibble: 2 x 2
  date      n
<date>   <int>
1 2011-05-09     1
2 2011-05-19   1513

```

```

si_au_prc_daily |>
  filter(dayssince == 40682) |>
  count(dateymd)

```

```
# A tibble: 1 x 2
  dateymd      n
  <date>    <int>
1 2011-05-19 1514
```

So in saving to parquet, we keep dateymd, but rename it to date for consistency across data sets.

```
si_au_prc_daily |>
  select(-date) |>
  rename(date = dateymd) |>
  write_parquet(sink = si_au_prc_daily_pq) |>
  system.time()
```

One issue with the code above is that it is quite slow and requires the full data set to be loaded in RAM. Given that si_au_prc_daily occupies 4.08 GB of RAM when loaded, this can be a problem if you have modest computing resources.

An alternative approach would be to use DuckDB's facility for reading CSV files and writing to parquet files. The small export_parquet() function accepts a remote data frame in a DuckDB connection and writes it to parquet.

```
export_parquet <- function(df, file) {
  db <- df[["src"]][["con"]]
  df <- dplyr::collapse(df)
  sql <- paste0("COPY (", dbplyr::remote_query(df),
                ") TO '", file, "'")
  DBI::dbExecute(db, sql)
  invisible(df)
}
```

The following code creates a DuckDB connection, then uses that connection to read the CSV file and then calls export_parquet() to write it the data to a parquet file. This is an order of magnitude faster than the read_csv() code above, yet seems to make no demands on RAM.

```
db <- dbConnect(duckdb::duckdb())

si_au_prc_daily <-
  tbl(db, str_c("read_csv('", si_au_prc_daily_csv, "',
                  DateFormat = '%Y%m%d',
                  types = {'dateymd': 'DATE',
                           'dayssince': 'INTEGER',
```

```

        'weekday': 'INTEGER',
        'monthend': 'BOOLEAN',
        'seniorsecurity': 'INTEGER'})"),
    name = "si_au_prc_daily") |>
select(-date) |>
rename(date = dateymd) |>
export_parquet(file = si_au_prc_daily_pq) |>
system_time()

```

```

user  system elapsed
29.060  1.342   6.246

```

3.4.1 Identifying the primary key

Obviously gcode and date are going to be part of any primary key, but we quickly deduce from the documentation supplied by SIRCA that a single gcode can be associated with multiple securities at one time and that seniorsecurity is used to distinguish these. This suggests (gcode, date, seniorsecurity) as a candidate primary key, so let's check this.

First, does each combination of (gcode, date, seniorsecurity) identify a single row?

```

si_au_prc_daily |>
  count(gcode, date, seniorsecurity, name = "num_rows") |>
  count(num_rows) |>
  collect()

```

```

# A tibble: 1 x 2
  num_rows      n
  <dbl>    <dbl>
1         1 8569251

```

Second, are there no NA values in the (gcode, date, seniorsecurity) combination?

```

si_au_prc_daily |>
  summarize(across(c(gcode, date, seniorsecurity),
                    \ (x) all(!is.na(x), na.rm = TRUE))) |>
  collect()

```

```

# A tibble: 1 x 3
  gcode date seniorsecurity
  <lgl> <lgl> <lgl>
1 TRUE  TRUE  TRUE

```

One thing to note here is that we checked the primary key using the DuckDB version of the data rather than the dplyr data frame (or tibble). One reason for this is that the code was much faster using the DuckDB version.

Now that we are done with our DuckDB connection, we can disconnect from it.

```
dbDisconnect(db)
```

3.5 The final script

I organize the code above (e.g., removed redundant elements) and placed it in a script [here](#). With the raw data in RAW_DATA_DIR and the necessary packages installed, I can create parquet data files by simply running the following code:¹⁹

```
Sys.setenv(RAW_DATA_DIR = "~/Library/CloudStorage/Dropbox/raw_data")
t <- tempdir()
Sys.setenv(DATA_DIR = t)
source("https://raw.githubusercontent.com/iangow/notes/main/import_sirca.R") |>
  system.time()
```

```
      user  system elapsed
28.899   1.145   6.373
```

4 Service-level agreement revisited

We return to our service-level agreement (SLA) to take stock of where we are after the above. Given that much of our focus above was on data types, I do not revisit that here and instead focus on those elements of the SLA, including those that I did not address above.

4.1 Storage format

We have chosen to use parquet files for our output. Table 4 provides some data on the parquet files we have produced for our hypothetical client (the *Understand* team). Assuming that the client is a group of colleagues at an institution with access to SIRCA, we (the *Curate* team) might just send a link to the Dropbox folder where we have stored the parquet files.

¹⁹Note that I set DATA_DIR to a different directory to avoid overwriting the files I just created and creating problems with Dropbox having to sync new files before it's even uploaded old ones.

Table 4: Data on processed parquet files

| file_name | size |
|---------------------------|-----------|
| si_au_prc_daily.parquet | 598.93 MB |
| si_au_ref_names.parquet | 830.48 kB |
| si_au_ref_trddays.parquet | 83.69 kB |
| si_au_retn_mkt.parquet | 422.42 kB |

4.2 Primary keys

Table 5 provides a summary of our analysis above of primary keys.

Table 5: SIRCA ASX EOD price collection: Primary keys

| Table | Primary key |
|-------------------|---------------------------------|
| si_au_ref_names | gcode, securityticker, listdate |
| si_au_prc_daily | gcode, date, seniorsecurity |
| si_au_retn_mkt | date |
| si_au_ref_trddays | date |

4.3 Good database principles

In general, I think one wants to be fairly conservative in considering database principles with a data library. If the data are workable and make sense in the form they come in, then it may make most sense to keep them in that form.

The SIRCA ASX EOD data are organized into four tables with easy-to-understand primary keys and a fairly natural structure. At some level, the two primary tables are `si_au_ref_names` and `si_au_prc_daily`.²⁰ These two tables are naturally distinct, with one about companies and the other about daily security returns.

While there might be merit in splitting `si_au_prc_daily` into separate tables to reduce its size, it is actually quite manageable in its current form.

4.4 No manual steps

There are no manual steps in creating the parquet files except for the initial download of the CSV files from SIRCA. While some data vendors allow users to download files using scripts

²⁰It seems possible that `si_au_retn_mkt` and `si_au_ref_trddays` are generated from `si_au_prc_daily`.

(e.g., the scripts I have [here](#) for WRDS), this does not appear to be an option for SIRCA. But once the data have been downloaded, the subsequent steps are automatic.

While some of the checks and data-cleaning had manual elements (e.g., identifying the near-duplicate with `Gcode=="rgwb1"` in `si_au_ref_names`), the resulting code implements the fix in an automated fashion. So long as the SIRCA data remain unchanged, the fix will continue to work.

4.5 Documentation

A important principle here is that the code for processing the data is documentation in its own right. Beyond that the document you are reading now is a form of documentation. If the goal of this document were to provide details explaining the process used to produce the final data sets, then it might make sense to edit this document to reflect that different purpose, but in many ways I hope this document already acts as good documentation.

4.6 Update process

In some ways, the update process is straightforward: when new CSV files become available, download them into `RAW_DATA_DIR` and run the script. However, it would probably be necessary to retrace some of the steps above to ensure that no data issues have crept in (e.g., duplicated keys). It may make sense to document the update process as part of performing it the first time.

4.7 Data version control

I achieve a modest level of data version control by using Dropbox, which offers the ability to restore previous versions of data files. As discussed earlier, version control of data is a knotty problem.

References

- Broman, Karl W., and Kara H. Woo. 2018. "Data Organization in Spreadsheets." *The American Statistician* 72 (1): 2–10. <https://doi.org/10.1080/00031305.2017.1375989>.
- Welch, Ivo. 2019. "Editorial: An Opinionated FAQ." *Critical Finance Review* 8 (1-2): 19–24. <https://doi.org/10.1561/104.00000077>.
- Wickham, Hadley, Mine Çetinkaya-Rundel, and Garrett Grolemund. 2023. *R for Data Science*. Sebastopol, CA: O'Reilly Media. <https://books.google.com/books?id=TiLEEAAQBAJ>.