# SIRCA ASX End of Day (EOD) collection

Ian D. Gow*

7 July 2024

## Overview

SIRCA's ASX EOD (end of day) collection provides daily prices for ASX-listed companies and facilitates reliable measurement of security returns, with all data being retained for delisted companies. This note, based on SIRCA's own `Guide to ASX End of Day Prices.pdf`, provides an introduction to the SIRCA ASX EOD collection.

The SIRCA ASX EOD collection includes the tables listed in Table 1. More details can be found in my separate document on importing SIRCA ASX EOF data here.

Table 1: SIRCA ASX EOD price collections

| Table | Description | Primary key |
|---|---|---|
| `si_au_ref_names` | Name and ticker histories for listed companies from January 2000 | `gcode`, `securityticker`, `listdate` |
| `si_au_prc_daily` | Complete daily price, volume and value histories | `gcode`, `date`, `seniorsecurity` |
| `si_au_retn_mkt` | Daily value- and equal-weighted whole market returns | `date` |
| `si_au_ref_trddays` | Record of ASX trading dates since January 2000 | `date` |

---

*University of Melbourne, ian.gow@unimelb.edu.au

All company names and ticker codes are recorded in a separate table—`si_au_ref_names`—which links names and tickers through time with a permanent "group code" identifier created by SIRCA named `gcode`. SIRCA designed `gcode` to allow users to build price and return series for series for a company's shares over time even if its ticker code changes. The same `gcodes` are used across a number of SIRCA's data sets.

he `si_au_prc_daily` table includes all end-of-day trade prices for the equity securities of all ASX companies starting from January 2000. This table also includes dividend and franking events; capital returns; adjustments for numerous corporate action events, such as splits, consolidation, bonus issues, renounceable and non-renounceable issues; and total daily traded volume and value; and, the number of issued shares.

Other components of the SIRCA ASX EOD library are `si_au_retn_mkt` which provides value- and equal-weighted all-of-market daily returns, which SIRCA generates from all observable daily company returns.

SIRCA provides `si_au_ref_trdday`, which identifies all ASX trading days since the start of January 2000 and can be used to identify gaps in company price series, from suspensions or thin trading.

Finally SIRCA provides a detailed description of all tables and fields in a data dictionary.

While this note is based on SIRCA's own `Guide to ASX End of Day Prices.pdf`, it goes beyond that guide in a number of respects. First, I provide detailed instructions on preparing the SIRCA ASX EOD data for use and illustrate analysis using DuckDB and parquet files, a high-performance state-of-the-art approach to data analysis. The SIRCA guide assumes that the user has access to an SQL database containing the four tables listed in Table 1, but provides no guidance on creating that database.

Second, I provide output—both tables and graphs—for the example queries provided here. While the SIRCA guide describes observed patterns, it does not provide the output from its queries.

Third, I expand on or refine the queries used in the SIRCA guide. For example, where the SIRCA guide suggests the use of `dayssince` variables to identify non-trading days, I propose a more robust approach.

The code in this note uses the packages listed below, plus the `duckdb` package.[1] This note was written using [Quarto](#) and compiled with [RStudio](#), an integrated development environment (IDE) for working with R. The source code for this note is available [here](#) and the latest PDF version is available [here](#).

---

[1] Execute `install.packages(c("tidyverse", "DBI", "duckdb", "arrow", "farr", "dbplyr"))` within R to install all the packages you need to run the code in this note. While `duckdb` and `arrow` are not listed below, they are needed to run the download script and to create the database we will use.

```r
library(tidyverse)
library(DBI)
library(dbplyr, warn.conflicts = FALSE)
library(farr)
```

## Getting SIRCA ASX EOF data

SIRCA supplies the data we will use as four compressed CSV files. The original SIRCA note assumes that you have processed the data into an SQL database, but does not provide any details for doing this. In contrast, I provide the code needed to prepare the data for the analysis below. I merely assume that you have access to SIRCA and have downloaded the raw data from SIRCA along the lines discussed below.[2] Once you have done that you should be able to execute all the code in this note and thereby produce all the output (tables and graphs) contained herein.

You should download these files and place them *as is* in a single `sirca` directory and you should edit the code below to tell my script where to look for that directory.

In my case, `sirca` is found in `~/Library/CloudStorage/Dropbox/raw_data` and so I execute the following command in RStudio:

```r
Sys.setenv(RAW_DATA_DIR = "~/Library/CloudStorage/Dropbox/raw_data")
```

Table 2 provides details on the contents of the `sirca` subdirectory of `RAW_DATA_DIR` on my computer.

Table 2: Data on supplied CSV files from SIRCA

| file_name | size |
|---|---|
| AuditFees-2024-08.csv.gz | 560.49 kB |
| MasterCompany-2024-08.csv.gz | 556.56 kB |
| si_au_prc_daily.csv.gz | 365.44 MB |
| si_au_ref_names.csv.gz | 591.81 kB |
| si_au_ref_trddays.csv.gz | 59.65 kB |
| si_au_retn_mkt.csv.gz | 352.44 kB |
| SIRCA EOD Data Dictionary.xlsx | 174.3 kB |

The script also needs to know where to put the processed data files, which will be in parquet format. You should edit the following line to refer to a location on *your* computer that you would like to use to store processed SIRCA data.

---

[2]See SIRCA's documentation for details on getting the data.

```r
Sys.setenv(DATA_DIR = "~/Library/CloudStorage/Dropbox/pq_data")
```

With the needed packages installed, and `RAW_DATA_DIR` and `DATA_DIR` set, we can run the script to process the SIRCA raw data with the following line. On my computer, this script takes about 6 seconds to run. If you wish to learn more about what the script is doing, please see the separate document on importing SIRCA ASX EOF data here.

```r
source("https://raw.githubusercontent.com/iangow/notes/main/import_sirca.R",
       echo = FALSE) |>
  system.time()
```

```
   user  system elapsed
 33.706   1.490  17.007
```

Details on the resulting parquet files that I have on my computer are provided in Table 3.

Table 3: Data on processed parquet files

| file_name | size |
| --- | --- |
| ann_fin.parquet | 78.75 MB |
| asic_short_interest.parquet | 27.39 MB |
| ma_anz_company.parquet | 843.87 kB |
| ma_anz_fdmt_auditfees.parquet | 595.9 kB |
| si_au_prc_daily.parquet | 598.93 MB |
| si_au_ref_names.parquet | 830.48 kB |
| si_au_ref_trddays.parquet | 83.69 kB |
| si_au_retn_mkt.parquet | 422.42 kB |

## Examples

The following examples demonstrate how to do certain analyses using the SIRCA ASX EOD data. These examples are based on examples provide in the SIRCA guide. Specifically, I show how one can use the SIRCA ASX EOD collection to do the following analyses:

1. Find a `gcode` from a company's name or ticker code
2. Apply the `cumulativefactor` field to adjust prices for different dividend and corporate action events
3. Generate and plot total shareholder returns
4. Use the `dayssince` column to identify return intervals between consecutive trades

5. Use the `seniorsecurity` column to focus on the residual risk security for each company
6. Handle negative factors and zero volumes
7. Calculate a cumulative factor excluding dividends
8. Segment trading activity by trade type and venue

All of the examples here use an in-memory DuckDB database connection. The following code creates this database connection and reads in three of the four SIRCA ASX EOD tables. In the cases of `si_au_ref_names`, the code also names that table so that we can refer to it using SQL written "by hand".[3]

```
db <- dbConnect(duckdb::duckdb())

si_au_ref_names <-
  load_parquet(db, "si_au_ref_names", "sirca") |>
  compute(name = "si_au_ref_names")

si_au_prc_daily <- load_parquet(db, "si_au_prc_daily", "sirca")
si_au_ref_trddays <- load_parquet(db, "si_au_ref_trddays", "sirca")
```

## 1. Finding `gcodes` from company names or ticker codes

One way of searching for a gcode is to look up the company name. For example, we could search for every gcode with a company name including `WESTPAC`. Because we specified `compute(name = "si_au_ref_names")`, we can refer to that table in SQL like the following:

```
SELECT gcode, seniorsecurity, securityticker, abbrevcompanyname
FROM si_au_ref_names
WHERE fullcompanyname LIKE '%WESTPAC%'
```

Table 4: Securities matching `WESTPAC`: SQL

| gcode | seniorsecurity | securityticker | abbrevcompanyname |
|-------|----------------|----------------|-------------------|
| wbc1  | 1              | WBC            | WESTPAC BANKING CORP |
| wot1  | 1              | WOT            | WESTPAC OFFICE TRUST |
| wot1  | 1              | WOTCA          | WESTPAC OFFICE TRUST |
| wpt1  | 1              | WPT            | WESTPAC PROP. TRUST |

However, the same query could be run using `tidyverse` code:

---

[3]While we access the data in a database throughout, most of the SQL is generated from `tidyverse` (R) code rather than being written by us directly.

```
si_au_ref_names |>
  filter(str_like(fullcompanyname, '%WESTPAC%')) |>
  select(gcode, seniorsecurity, securityticker, abbrevcompanyname) |>
  collect()
```

Table 5: Securities matching `WESTPAC`: Tidyverse

| gcode | seniorsecurity | securityticker | abbrevcompanyname |
|-------|----------------|----------------|-------------------|
| wbc1 | 1 | WBC | WESTPAC BANKING CORP |
| wot1 | 1 | WOT | WESTPAC OFFICE TRUST |
| wot1 | 1 | WOTCA | WESTPAC OFFICE TRUST |
| wpt1 | 1 | WPT | WESTPAC PROP. TRUST |

Behind the scenes, the `tidyverse` package is translating our code into SQL:

```
si_au_ref_names |>
  filter(str_like(fullcompanyname, '%WESTPAC%')) |>
  select(gcode, seniorsecurity, securityticker, abbrevcompanyname) |>
  show_query()
```

```
<SQL>
SELECT gcode, seniorsecurity, securityticker, abbrevcompanyname
FROM si_au_ref_names
WHERE (fullcompanyname LIKE '%WESTPAC%')
```

Alternatively, one can also search by ticker code, as seen in Table 6 using the `securityticker` of `ANZ`.

```
SELECT gcode, securityticker, abbrevcompanyname
FROM si_au_ref_names
WHERE securityticker = 'ANZ'
```

Table 6: Securities with ticker `ANZ`: SQL

| gcode | securityticker | abbrevcompanyname |
|-------|----------------|-------------------|
| anz1 | ANZ | AUSTRALIA AND NZ |
| anz1 | ANZ | ANZ GROUP HOLDINGS |

Again the same query could be run using `tidyverse` code, with results show in Table 7. Because it is so straightforward to run SQL queries using R (`tidyverse`) code, we will just provide R code going forward.[4] Note that using R code greatly facilitates bringing the data into R for analysis or (as we will do here) data visualization.

```
si_au_ref_names |>
  filter(securityticker == 'ANZ') |>
  select(gcode, seniorsecurity, securityticker, abbrevcompanyname) |>
  collect()
```

Table 7: Securities with ticker `ANZ`: Tidyverse

| gcode | seniorsecurity | securityticker | abbrevcompanyname |
|-------|----------------|----------------|-------------------|
| anz1 | 1 | ANZ | AUSTRALIA AND NZ |
| anz1 | 1 | ANZ | ANZ GROUP HOLDINGS |

As another example, suppose one is interested in Arena REIT which has a `companyticker` of `ARF`. Searching for this ticker code reveals the `gcode` of Arena REIT is `arf2`. This `gcode` can then be used to search the `si_au_prc_daily` table for information about the securities of Arena REIT. Results of this search are seen in Table 8.

```
si_au_ref_names |>
  filter(companyticker == 'ARF', seniorsecurity == 1L) |>
  select(gcode, securityticker, abbrevcompanyname,
         listdate, delistdate) |>
  arrange(listdate) |>
  collect()
```

Table 8: Securities with ticker `ARF`

| gcode | securityticker | abbrevcompanyname | listdate | delistdate |
|-------|----------------|-------------------|----------|------------|
| arf1 | ARF | ARROWFIELD GROUP | 1987-11-12 | 2000-09-20 |
| arf2 | ARF | ARENA REIT | 2013-06-13 | 2013-12-10 |
| arf2 | ARFDA | ARENA REIT | 2013-12-11 | 2013-12-19 |
| arf2 | ARF | ARENA GROUP | 2013-12-20 | 2014-01-19 |
| arf2 | ARF | ARENA REIT | 2014-01-20 | 2014-12-08 |
| arf2 | ARFDC | ARENA REIT | 2014-12-09 | 2014-12-15 |
| arf2 | ARF | ARENA REIT | 2014-12-16 | NA |

---

[4]Users with a background in SQL may find the SQL primer I wrote [here] to be a useful introduction to the `dplyr` package (this is the component of the `tidyverse` package that provides the relevant functions).

The previous search reveals that in 2000, the ticker code ARF was then associated with Arrow-field Group Limited. Arrowfield Group Limited is a different entity to Arena REIT, which was listed in 2013, so the two entities have separate gcodes.

A search for `companyticker` of ARF also shows that in 2013, the `securityticker` of Arena REIT briefly changed from ARF to ARFDA, and then back to ARF, due to conversions to and from deferred units. If the holders of units in Arena REIT in 2013 participated in these conversions, then it makes sense to consider them as a single security, which is possible if we accumulate returns using the `gcode` of `arf2`, which remains unchanged throughout this period.

## 2. Adjusting for the effects of corporate actions

Figure 1 shows a large drop in the closing price of BHP in late June 2001.

```
si_au_prc_daily |>
  filter(gcode == 'bhp1', seniorsecurity == 1L,
         between(date, '2001-01-01', '2001-12-31')) |>
  ggplot(aes(x = date, y = close)) +
  geom_line()
```



Figure 1: BHP stock price during 2001

8

Examining the `coraxdescription` column, it seems likely that the change is due to a `1:0.94` bonus issue.

```
si_au_prc_daily |>
  filter(gcode == 'bhp1', seniorsecurity == 1L,
         between(date, '2001-01-01', '2001-12-31')) |>
  filter(!is.na(coraxdescription)) |>
  select(gcode, date, close, coraxdescription) |>
  collect()
```

Table 9: Corporate action events for BHP in 2001

| gcode | date | close | coraxdescription |
|---|---|---|---|
| bhp1 | 2001-06-29 | 10.39 | 1:0.94 bonus issue |

The `coraxdescription` column provides details of corporate action (CORAX) events, when available. The `numberofdilutionevents` field will always show a value greater than 0 when a dilution event (CORAX or dividend) has occurred and `numberofcoraxevents > 0` indicates CORAX events, even if `coraxdescription` is not available. Likewise, `numberofdividendevents > 0` can be used to find all dividend events, even when data are not available in other descriptive fields .

```
si_au_prc_daily |>
  filter(gcode == 'bhp1', seniorsecurity == 1L,
         between(date, '2001-06-27', '2001-07-02')) |>
  select(gcode, date, close, factor, numberofcoraxevents) |>
  collect()
```

Table 10: Values of `factor` for BHP around 27 June 2001

| gcode | date | close | factor | numberofcoraxevents |
|---|---|---|---|---|
| bhp1 | 2001-06-27 | 20.946 | 1.000 | 0 |
| bhp1 | 2001-06-28 | 21.420 | 1.000 | 0 |
| bhp1 | 2001-06-29 | 10.390 | 0.484 | 1 |
| bhp1 | 2001-07-02 | 10.480 | 1.000 | 0 |

How is `factor` calculated? As seen in Table 10, on most dates `factor` will be 1, but when a CORAX or dividend event occurs, `factor` will reflect the factor that allows the previous day's price to be compared with the current one. With a `1:0.94` bonus issue, if I have 0.94 shares one day, I will have 1.94 shares the next, so `factor` equals 0.94/1.94 = 0.485. In other words, the

share price of 21.420 on 28 June 2001 is equivalent to a share price of 21.420 × 0.485 = 10.379 on 29 June 2001.

The variable `cumulativefactor` is calculated from `factor` to facilitate adjustment of prices along the whole time series. In creating `cum_factor_calcs`, I replicate the calculation of `cumulativefactor` from `factor`. Note that the calculation of `cumulativefactor` moves from the *end* of the price series (implied by `window_order(desc(date))`) for each security (implied by `group_by(gcode, seniorsecurity)`) and accumulates the absolute value of `factor` in a multiplicative fashion.[5] The calculation uses the `lag()` function because the first date we want to apply `factor` for 29 June 2001 to prices is the "next" date (in the reverse-ordered price series) or 28 June 2001.[6]

```
cum_factor_calcs <-
  si_au_prc_daily |>
  group_by(gcode, seniorsecurity) |>
  window_order(desc(date)) |>
  mutate(cum_factor_calc = exp(cumsum(log(abs(factor))))) |>
  mutate(cum_factor_calc = lag(cum_factor_calc) * sign(lag(factor))) |>
  window_order() |>
  ungroup() |>
  select(gcode, seniorsecurity, date, close, factor,
         cumulativefactor, cum_factor_calc)
```

Table 11 presents `cumulativefactor`, as calculated by SIRCA, and `cum_factor_calc`, where I replicate the calculation of `cumulativefactor` from `factor`.

```
cum_factor_calcs |>
  filter(gcode == 'bhp1', seniorsecurity == 1L,
         between(date, '2001-06-27', '2001-07-02')) |>
  select(gcode, date, close, factor, cumulativefactor, cum_factor_calc) |>
  arrange(date) |>
  collect()
```

Table 11: Calculated values of `factor` for BHP around 27 June 2001

| gcode | date | close | factor | cumulativefactor | cum_factor_calc |
|-------|------|-------|--------|------------------|-----------------|
| bhp1 | 2001-06-27 | 20.946 | 1.000 | 0.153 | 0.153 |
| bhp1 | 2001-06-28 | 21.420 | 1.000 | 0.153 | 0.153 |
| bhp1 | 2001-06-29 | 10.390 | 0.484 | 0.315 | 0.315 |

---

[5]The reason for taking absolute values and accounting for the sign of `factor` is discussed below in Section .

[6]Note that the calculation of `cum_factor_calc` occurs on two separate lines, as DuckDB does not allow nesting of **window functions**, such as `cumsum()` and `lag()`. For more on window functions, see here.

Table 11: Calculated values of `factor` for BHP around 27 June 2001

| gcode | date | close | factor | cumulativefactor | cum_factor_calc |
|-------|------|------|--------|------------------|-----------------|
| bhp1 | 2001-07-02 | 10.480 | 1.000 | 0.315 | 0.315 |

The `cumulativefactor` column can be used to adjust the closing price for the effects of corporate actions, such stock splits or entitlement offers, and dividends. Simply multiplying the `close` column by the `cumulativefactor` column will produce the adjusted price.

```
si_au_prc_daily |>
  filter(gcode == 'bhp1', seniorsecurity == 1L,
         between(date, '2001-01-01', '2001-12-31')) |>
  mutate(adjustedprice = close * cumulativefactor) |>
  pivot_longer(c(adjustedprice, close),
               names_to = "variable", values_to = "price") |>
  ggplot(aes(x = date, y = price, color = variable)) +
  geom_line()
```



Figure 2: BHP adjusted stock price during 2001

In Figure 2, `adjustedprice` series is everywhere lower than `close` because `cumulativefactor` adjusts for all subsequent CORAX events and these tend to cause adjusted prices to be lower as

11

one moves back through time (e.g., bonus issues or dividends). The important thing is that the resulting `adjustedprice` series is consistent over its entire history and can be used to reliably measure returns for `bhp1` between any two trading dates.

Exactly the same process for `cumulativefactor` applies for dividends as well as corporate actions. `AAA` (gcode: `aaa2`) is an exchange-traded fund that deposits money in accounts with Australian banks and pays regular dividends. The effect of its dividends on its closing price can be observed in Figure 3.

```
si_au_prc_daily |>
  filter(gcode == 'aaa2', seniorsecurity == 1L,
         between(date, '2017-01-01', '2018-12-31')) |>
  mutate(adjustedprice = close * cumulativefactor) |>
  select(date, close, adjustedprice) |>
  pivot_longer(-date, names_to = "variable", values_to = "price") |>
  ggplot(aes(x = date, y = price, color = variable)) +
  geom_line()
```



Figure 3: Adjusted and unadjusted closing prices for `AAA`

### 3. Plotting a distribution of price relatives for a security

The following code calculates `prel`, the price relative or gross shareholder return, for securities on `si_au_prc_daily`. By using `cumulativefactor`, it adjusts for corporate actions and dividends.

```
prels <-
  si_au_prc_daily |>
  mutate(adjustedprice = close * cumulativefactor) |>
  group_by(gcode, seniorsecurity) |>
  window_order(date) |>
  mutate(prel = adjustedprice / lag(adjustedprice)) |>
  ungroup() |>
  window_order()
```

Figure 4 shows the distribution of returns from Commonwealth Bank.

```
prels |>
  filter(!is.na(prel)) |>
  filter(gcode == 'cba1', seniorsecurity == 1L) |>
  ggplot(aes(x = prel)) +
  geom_histogram(binwidth = 0.005)
```



Figure 4: Distribution of daily returns for Commonwealth Bank

## 4. `dayssince` column

It is important to note that price relatives calculated in the previous section may not always relate to consecutive trading days. The following code calculates the number of days between consecutive observations for a given security on `si_au_prc_daily`.

```
elapsed_days <-
  si_au_prc_daily |>
  group_by(gcode, seniorsecurity) |>
  window_order(date) |>
  mutate(days_elapsed = dayssince - lag(dayssince))
```

Figure 5 shows the distribution of `days_elapsed`, the number of elapsed days between trading dates calculated using using the `datesince` column, for Commonwealth Bank. Although CBA is a stock that is consistently traded, a less-liquid security may show large gaps in trading activity, leading to price relatives that span longer time periods.

```
elapsed_days |>
  count(days_elapsed, sort = TRUE) |>
  filter(gcode == 'cba1', seniorsecurity == 1L) |>
  filter(!is.na(days_elapsed)) |>
  ggplot(aes(x = days_elapsed, y = n)) +
  geom_col() +
  scale_x_continuous(breaks = 1:5)
```
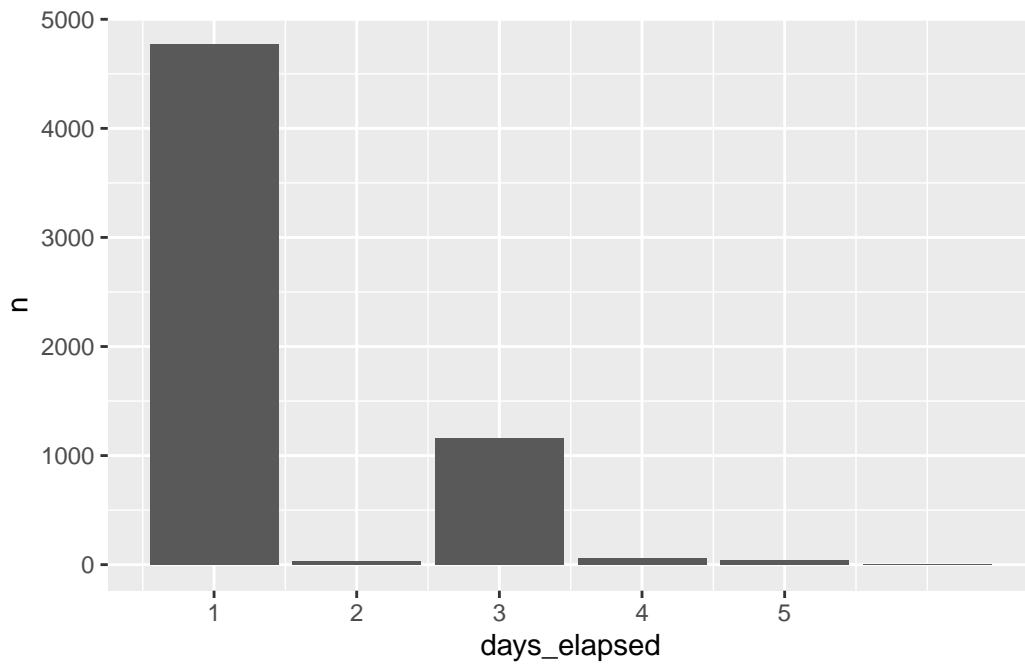
Figure 5: Distribution of days between trading dates for Commonwealth Bank

As a measure of the liquidity of a security, `days_elapsed` is problematic because it does not distinguish between days on which the market is open and those on which it is closed. We can improve on this measure using data from `si_au_ref_trddays`, a sample of which is shown in Table 12.

```
si_au_ref_trddays |> collect(n = 10)
```

Table 12: Sample rows from `si_au_ref_trddays`

| date | dayssince | weekday | count |
|---|---|---|---|
| 2000-01-04 | 36529 | 3 | 936 |
| 2000-01-05 | 36530 | 4 | 958 |
| 2000-01-06 | 36531 | 5 | 949 |
| 2000-01-07 | 36532 | 6 | 927 |
| 2000-01-10 | 36535 | 2 | 961 |
| 2000-01-11 | 36536 | 3 | 984 |
| 2000-01-12 | 36537 | 4 | 968 |
| 2000-01-13 | 36538 | 5 | 971 |
| 2000-01-14 | 36539 | 6 | 966 |
| 2000-01-17 | 36542 | 2 | 976 |

Using an approach described in more detail here, in place of `dayssince`, we can create a variable `td` to represent the "trading date" for each date on `si_au_ref_trddays` where `td` equals 1 on the first trading date, 2 the second trading date, and so on.

```
trading_days <-
  si_au_ref_trddays |>
  window_order(date) |>
  mutate(td = row_number()) |>
  distinct(date, td) |>
  arrange(date) |>
  compute()
```

With `trading_days` in hand, we can calculate `tdays_elapsed` as the number of trading dates between the current date and the previous date on `si_au_prc_daily` for each security and date.

```
tdays_elapsed_df <-
  si_au_prc_daily |>
  inner_join(trading_days, by = "date") |>
  group_by(gcode, seniorsecurity) |>
  window_order(date) |>
  mutate(tdays_elapsed = td - lag(td),
         lag_date = lag(date)) |>
  select(gcode, seniorsecurity, date, lag_date, tdays_elapsed) |>
  ungroup()
```

Table 13 provides data on our improved measure of trading days between trading dates for both Commonwealth Bank (`cba1`) and a less liquid security (`1st1`). In Table 13, it can be seen that there are very few cases in which the trading days between dates is more than one for `cba1`, but quite a few such cases for `1st1`.

```
tdays_elapsed_df |>
  filter(gcode %in% c('cba1', '1st1'), seniorsecurity == 1L,
         !is.na(tdays_elapsed)) |>
  count(gcode, tdays_elapsed) |>
  pivot_wider(names_from = "gcode", values_from = "n", values_fill = 0) |>
  arrange(tdays_elapsed) |>
  collect()
```

Table 13: Trading days between trading dates: `cba1` and `1st1`

| tdays_elapsed | cba1 | 1st1 |
|--------------:|-----:|-----:|
| 1 | 6060 | 1228 |
| 2 | 2 | 150 |
| 3 | 1 | 67 |
| 4 | 1 | 27 |
| 5 | 0 | 13 |
| 6 | 0 | 6 |
| 7 | 0 | 2 |
| 8 | 0 | 5 |
| 9 | 0 | 1 |
| 10 | 0 | 1 |
| 11 | 0 | 2 |
| 134 | 0 | 1 |

Table 14 provides additional information on the apparent gaps in trading for Commonwealth Bank. We can use these data to investigate the cause of these gaps. Looking at the longest gap, it turns out there was a trading halt placed on 12 August 2015.

```
tdays_elapsed_df |>
  filter(gcode == 'cba1', seniorsecurity == 1L, tdays_elapsed > 1) |>
  collect()
```

Table 14: Gaps in trading for Commonwealth Bank

| gcode | seniorsecurity | date | lag_date | tdays_elapsed |
|-------|---------------:|------------|------------|--------------:|
| cba1 | 1 | 2000-03-10 | 2000-03-07 | 3 |
| cba1 | 1 | 2008-10-09 | 2008-10-07 | 2 |
| cba1 | 1 | 2015-08-17 | 2015-08-11 | 4 |
| cba1 | 1 | 2015-09-14 | 2015-09-10 | 2 |

A natural question might be whether there are dates on `si_au_prc_daily` not found on `si_au_ref_trddays`. Table 15 shows that there are, but a small number of securities (in most cases, just one) have data on `si_au_prc_daily` on those days. I leave it as an exercise for the reader to understand what's going on in these cases.

```
si_au_prc_daily |>
  distinct(date) |>
  anti_join(si_au_ref_trddays, by = "date") |>
```

```
inner_join(si_au_prc_daily, by = "date") |>
count(date) |>
mutate(wday = wday(date, label = TRUE)) |>
arrange(desc(n)) |>
collect(n = 10)
```

Table 15: Observations on `si_au_prc_daily` on non-trading days

| date | n | wday |
|------|---|------|
| 2014-01-01 | 3 | Wed |
| 2000-01-01 | 3 | Sat |
| 2014-06-01 | 2 | Sun |
| 2009-01-01 | 2 | Thu |
| 2014-12-20 | 1 | Sat |
| 2008-11-29 | 1 | Sat |
| 2010-02-27 | 1 | Sat |
| 2011-09-10 | 1 | Sat |
| 2002-06-08 | 1 | Sat |
| 2011-08-21 | 1 | Sun |

## 5. Using the `seniorsecurity` column

At times, some gcodes have multiple securities trading simultaneously and SIRCA provides the `seniorsecurity` field to distinguish different securities for a given firm. In Table 16, two classes of security are shown to be simultaneously trading for Telstra Corporation Ltd, whose gcode is `tls1`. These are evident from the different `securityticker` values: TLS and TLSCA.

```
si_au_prc_daily |>
  filter(gcode == 'tls1',
         between(date, "2008-05-01", "2008-05-07")) |>
  select(gcode, seniorsecurity, date, securityticker) |>
  arrange(date) |>
  collect()
```

Table 16: Sample of `securityticker` values for Telstra (`tls1`)

| gcode | seniorsecurity | date | securityticker |
|-------|----------------|------|----------------|
| tls1 | 0 | 2008-05-01 | TLSCA |
| tls1 | 1 | 2008-05-01 | TLS |

Table 16: Sample of `securityticker` values for Telstra (`tls1`)

| gcode | seniorsecurity | date | securityticker |
|-------|---------------:|------|----------------|
| tls1 | 0 | 2008-05-02 | TLSCA |
| tls1 | 1 | 2008-05-02 | TLS |
| tls1 | 0 | 2008-05-05 | TLSCA |
| tls1 | 1 | 2008-05-05 | TLS |
| tls1 | 0 | 2008-05-06 | TLSCA |
| tls1 | 1 | 2008-05-06 | TLS |
| tls1 | 0 | 2008-05-07 | TLSCA |
| tls1 | 1 | 2008-05-07 | TLS |

## 6. Negative `factor` values and zero `volumeonmkt` values

When there is either no observed trade price before an event or no price after the event, a `factor` of `-1` is assigned to that event. This can occur both in the beginning and the end of the lifetime of the security. Table 17 shows that relatively few observations have negative `factor` values.

```
si_au_prc_daily |>
  mutate(neg_factor = factor < 0) |>
  count(neg_factor)
```

Table 17: Number of observations by negative factors

| neg_factor | n |
|------------|--------:|
| FALSE | 8569038 |
| TRUE | 213 |

The calculation of `cumulativefactor` when `factor` is negative seems to follow the calculation of `cum_factor_calc` provided above. That is the absolute value is accumulated and multiplied by the sign of the applicable `factor` value. Table 18 shows the alignment of `cumulativefactor` and `cum_factor_calc` calculated in this way for a stock with a negative value of `factor`.

```
cum_factor_calcs |>
  filter(gcode == "par1", seniorsecurity == 1L,
         between(date, "2018-12-03", "2019-02-18")) |>
  select(-seniorsecurity) |>
  arrange(date)
```

Table 18: Calculating `cumulativefactor` with negative `factor` values

| gcode | date | close | factor | cumulativefactor | cum_factor_calc |
|-------|------|-------|--------|------------------|-----------------|
| par1 | 2018-12-11 | 0.003 | 1.000 | 0.765 | 0.765 |
| par1 | 2018-12-12 | 0.003 | 1.000 | 0.765 | 0.765 |
| par1 | 2018-12-14 | 0.004 | 1.000 | -0.765 | -0.765 |
| par1 | 2019-01-08 | NA | -0.765 | 1.000 | 1.000 |
| par1 | 2019-01-24 | 0.005 | 1.000 | 1.000 | 1.000 |
| par1 | 2019-02-18 | 0.009 | 1.000 | 1.000 | 1.000 |

Table 19 shows that the calculation used to produce `cum_factor_calc` does not always match the valye in `cumulativefactor`. In this case, it seems that `cumulativefactor` is mysteriously "reset" to 1 on 27 September 2018. Further research would be needed to determine

```
cum_factor_calcs |>
  filter(gcode == "gcm2", seniorsecurity == 1,
         factor != 1) |>
  select(-seniorsecurity) |>
  arrange(date)
```

Table 19: Case of mysterious `factor` values

| gcode | date | close | factor | cumulativefactor | cum_factor_calc |
|-------|------|-------|--------|------------------|-----------------|
| gcm2 | 2016-09-29 | NA | 0.990 | 0.917 | 0.879 |
| gcm2 | 2016-12-29 | 0.95 | 0.990 | 0.927 | 0.888 |
| gcm2 | 2017-03-30 | NA | 0.990 | 0.936 | 0.897 |
| gcm2 | 2017-06-29 | NA | 0.989 | 0.947 | 0.908 |
| gcm2 | 2017-09-28 | NA | 0.989 | 0.958 | 0.918 |
| gcm2 | 2017-12-28 | NA | 0.991 | 0.967 | 0.926 |
| gcm2 | 2018-03-28 | NA | 0.991 | 0.975 | 0.935 |
| gcm2 | 2018-06-28 | NA | 0.983 | 0.993 | 0.951 |
| gcm2 | 2018-09-27 | NA | 0.993 | 1.000 | 0.958 |
| gcm2 | 2018-12-28 | NA | -1.000 | 0.958 | 0.958 |
| gcm2 | 2019-03-28 | NA | 0.992 | 0.965 | 0.965 |
| gcm2 | 2019-06-06 | NA | 0.978 | 0.988 | 0.988 |
| gcm2 | 2019-06-27 | 1.79 | 0.993 | 0.994 | 0.994 |
| gcm2 | 2019-09-27 | NA | 0.994 | 1.000 | 1.000 |
| gcm2 | 2019-11-20 | NA | -1.000 | -1.000 | NA |

Table 20 flags other difficult-to-explain `cumulativefactor` values (excluding those where there are sign differences between `cum_factor_calc` and `cumulativefactor`). While further research would be needed to understand these, these are fortunately quite rare.

```
cum_factor_calcs |>
  filter(gcode != "gcm2") |>
  filter(abs(abs(cumulativefactor) - abs(cum_factor_calc)) > 0.001) |>
  distinct(gcode, cumulativefactor, cum_factor_calc) |>
  arrange(gcode, cumulativefactor) |>
  collect()
```

Table 20: Other difficult-to-explain `factor` values

| gcode | cumulativefactor | cum_factor_calc |
|-------|-----------------|-----------------|
| bsc1  | 4.000           | -1.000          |
| ikw1  | -40.471         | -4047.087       |
| ikw1  | 40.471          | 4047.087        |
| ikw1  | 50.479          | 5047.868        |
| npx1  | -1.000          | -0.689          |

The example provided in Table 21 shows dividends between 2000-03-06 and 2001-09-28 without any trading. As no trading was observed prior to these dividend events, the `factor` and `dividendfactor` fields contain a value of -1. This makes sense, as one could not meaningfully push the sequence of stock returns back to dates before 2002-03-15, as there are no traded prices. There is a non-negative factor value for 2002-03-18, presumably because there are prices reported after 2002-03-18.

```
si_au_prc_daily |>
  filter(gcode == 'npx1', date <= '2002-09-23', seniorsecurity == 1) |>
  select(gcode, date, close, dividend, factor, dividendfactor, volumeonmkt) |>
  collect()
```

Table 21: Negative `factor` values: npx1

| gcode | date       | close | dividend | factor | dividendfactor | volumeonmkt |
|-------|------------|-------|----------|--------|----------------|-------------|
| npx1  | 2000-03-06 | NA    | 0.065    | -1.000 | -1.000         | 0           |
| npx1  | 2000-09-29 | NA    | 0.054    | -1.000 | -1.000         | 0           |
| npx1  | 2001-03-19 | NA    | 0.068    | -1.000 | -1.000         | 0           |
| npx1  | 2001-09-28 | NA    | 0.059    | -1.000 | -1.000         | 0           |
| npx1  | 2002-03-15 | 3.34  | NA       | 1.000  | 1.000          | 500         |
| npx1  | 2002-03-18 | NA    | 0.067    | 0.978  | 0.978          | 0           |
| npx1  | 2002-07-17 | 2.95  | NA       | 1.000  | 1.000          | 200         |
| npx1  | 2002-09-23 | NA    | 0.078    | 1.000  | 1.000          | 0           |

Table 22 shows another example with a dividend on 2004-07-05. However, no price was observed after the event, and hence the `factor` and `dividendfactor` fields contain a value of -1. Note that there is a price in the `close` field on 2004-07-05 but it was not observed that day, after the dividend event. This is evident from the 0 value for `VolumeOnMkt`, and confirmed by `NA` or `0` values for `open`, `high`, `low`. This price is simply the previous observed trade price carried forward. This makes sense, as one could not meaningfully push the sequence of stock returns forward to dates after 2004-07-02, as there are no traded prices.

```
si_au_prc_daily |>
 filter(gcode == 'wsf1', date >= '2004-07-01', seniorsecurity == 1) |>
  select(gcode, date, close, dividend, factor, dividendfactor, volumeonmkt) |>
  collect()
```

Table 22: Dividend example: `npx1`

| gcode | date | close | dividend | factor | dividendfactor | volumeonmkt |
|---|---|---|---|---|---|---|
| wsf1 | 2004-07-01 | 15.48 | NA | 1 | 1 | 8968726 |
| wsf1 | 2004-07-02 | 15.60 | NA | 1 | 1 | 16453850 |
| wsf1 | 2004-07-05 | 15.60 | 0.136 | -1 | -1 | 0 |

```
si_au_prc_daily |>
  filter(adjustmentfactor < 0) |>
  count(adjustmentfactor)
```

Table 23: Frequency of negative `adjustmentfactor` values

| adjustmentfactor | n |
|---|---|
| -40.471 | 1 |
| -0.765 | 1 |
| -1.000 | 57 |

## 7. Calculating a cumulative factor excluding dividends

The provided `cumulativefactor` field of `si_au_prc_daily` is calculated by cumulating the `factor` column, whcih adjusts for both corporate actions and dividends. The following example shows how to calculate an adjustment *excluding* dividends, It uses the `adjustmentfactor` field, which provides dilution factors for just the CORAX events (when followed at some time by a valid `close` price).

Use the `adjustmentfactor` field, which does not account for dividends. Visualise the new adjustment and compare to the adjustment from the example in part 1 Note that the

`CorpAdjustedPrice`, which is calculated without including dividends, looks identical to the `close` price as no corporate actions have occurred within this time frame.

The following shows the effect that dividends can have on the adjusted price series. The `AdjustedPrice` series incorporates both CORAX factors and dividend factors, whereas the `CorpAdjustedPrice` series incorporates only CORAX adjustments and ignores dividends. The CORAX-only price series shows a visible fall at the time when the dividend occurs as the value of the dividend is not accounted for.

```
adj_rets <-
  si_au_prc_daily |>
  group_by(gcode, seniorsecurity) |>
  window_order(desc(date)) |>
  mutate(corporatefactor = exp(cumsum(log(abs(adjustmentfactor))))) |>
  mutate(corporatefactor = lag(corporatefactor) *
           sign(lag(adjustmentfactor))) |>
  mutate(CorpAdjustedPrice = corporatefactor * close,
         AdjustedPrice = close * cumulativefactor) |>
  window_order() |>
  ungroup()
```

To plot some date, we first construct `anz_cum`, which is a version of `adj_rets` focused on ANZ's stock price.

```
anz_cum <-
  adj_rets |>
  filter(gcode == "anz1", seniorsecurity == 1L) |>
  select(gcode, date, CorpAdjustedPrice, AdjustedPrice, close) |>
  compute()
```

Figure 6 provides a plot of `CorpAdjustedPrice` (no adjustment for dividends) and `AdjustedPrice` (adjusted for dividends) for ANZ. One thing that makes this plot difficult to interpret is that the adjustment factors are calculated retrospectively. So *going back in time* the plots "start" at the same point and "end" at different prices.

```
anz_cum |>
  pivot_longer(cols = ends_with("Price"),
               names_to = "series", values_to = "price") |>
  filter(!is.na(price)) |>
  ggplot(aes(x = date, y = price, color = series)) +
  geom_line() +
  theme(legend.position = "top")
```

Figure 6: ANZ: `CorpAdjustedPrice` and `AdjustedPrice`

Figure 7 provides a more intuitive presentation. Rather than calculating returns using the adjusted prices directly, an alternative measure of returns is constructed by calculating a price relative using adjusted prices and then accumulating those returns. As can be seen in Figure 7, the two price series start at the same point (no scare quotes because this plot is going forward in time) and end at different points. As would be expected the cumulative returns without dividends are significantly lower by the end of the price series.

```
anz_cum |>
  group_by(gcode) |>
  window_order(date) |>
  mutate(across(c(AdjustedPrice, CorpAdjustedPrice),
                \(x) coalesce(x / lag(x), 1)),
         across(c(AdjustedPrice, CorpAdjustedPrice),
                \(x) exp(cumsum(log(x))))) |>
  window_order() |>
  pivot_longer(cols = ends_with("Price"),
               names_to = "series", values_to = "price") |>
  filter(!is.na(price)) |>
  ggplot(aes(x = date, y = price, color = series)) +
  geom_line() +
  theme(legend.position = "top")
```

Figure 7: ANZ: Fixed plot of `CorpAdjustedPrice` and `AdjustedPrice`

Table 24 presents an example of an even larger difference between the different forms of adjusted price. The dividend of 0.4497 on `2012-07-09` precedes a fall in `close` price from 0.575 to 0.019. Adjusting only for CORAX events clearly leads to significantly different measures of share price performance when dividends are also present. If done with case, CORAX adjustments might also be used to standardise earnings information through time.

```
adj_rets |>
 filter(gcode == "dmg1",
        between(date, "2012-07-01", "2012-07-13"),
        seniorsecurity == 1) |>
  select(gcode, date, close, dividend, CorpAdjustedPrice, AdjustedPrice) |>
  arrange(date) |>
  collect()
```

Table 24: Something about `dmg1`

| gcode | date | close | dividend | CorpAdjustedPrice | AdjustedPrice |
|-------|------|-------|----------|-------------------|---------------|
| dmg1 | 2012-07-02 | 0.570 | NA | 0.071 | 0.018 |
| dmg1 | 2012-07-03 | 0.570 | NA | 0.071 | 0.018 |
| dmg1 | 2012-07-04 | 0.575 | NA | 0.072 | 0.018 |

Table 24: Something about `dmg1`

| gcode | date | close | dividend | CorpAdjustedPrice | AdjustedPrice |
|---|---|---|---|---|---|
| dmg1 | 2012-07-05 | 0.575 | NA | 0.072 | 0.018 |
| dmg1 | 2012-07-06 | 0.575 | NA | 0.072 | 0.018 |
| dmg1 | 2012-07-09 | 0.019 | 0.45 | 0.019 | 0.019 |
| dmg1 | 2012-07-10 | 0.020 | NA | 0.020 | 0.020 |
| dmg1 | 2012-07-11 | 0.019 | NA | 0.019 | 0.019 |
| dmg1 | 2012-07-12 | 0.018 | NA | 0.018 | 0.018 |
| dmg1 | 2012-07-13 | 0.018 | NA | 0.018 | 0.018 |

## 8. Segmentation by trade type

The `si_au_prc_daily` table also contains information on the count, volume, and value of trades by various categories. This section provides examples of aggregating trading activity by different trade types:

1. Trading activity across the whole market
2. Segmentation by on- versus off-market trades
3. Proportion of on-market non-crossing trades that are carried out through ASX Centre Point
4. Comparison of lit-pool and dark-pool trading
5. Proportion of dark market trades that are carried out through ASX Centre Point

Figure 8 shows the value of trading activity across the year plotted against time. Trading activity can vary significantly from month to month.

```
si_au_prc_daily |>
  mutate(month = floor_date(date, "month")) |>
  group_by(month) |>
  summarize(ValueWholeMkt = sum(valueonmkt + valueoffmkt, na.rm = TRUE)) |>
  ggplot(aes(x = month, y = ValueWholeMkt)) +
  geom_line()
```

Figure 8: Trading activity

```
si_au_prc_daily |>
  mutate(month = floor_date(date, "month")) |>
  group_by(month) |>
  summarize(across(c(valueonmkt, valueoffmkt), \(x) sum(x, na.rm = TRUE))) |>
  pivot_longer(-month, names_to = "location", values_to = "value") |>
  ggplot(aes(x = month, y = value, color = location)) +
  geom_line() +
  theme(legend.position = "top")
```

Figure 9: Trading activity: On- versus off-market

As mentioned above, it is possible to segment the market by the visibility of trades. In the lit market, the order book is public and all orders (bid and offer) are visible to all participants. In contrast, in the dark market, the order book is not visible until trades are executed. The dark pool consists of both on-market and off-market crossing trades, as well as any Centre Point trades. This following section shows the distribution of activity across the lit and dark markets over time. Note: As our Centre Point trade measures include crossing trades, Centre Point crossing trade volumes need to be subtracted to avoid double-counting these trades in the calculation of the dark pool trading.

```
si_au_prc_daily |>
  mutate(month = floor_date(date, "month")) |>
  filter(volumeonmkt > 0) |>
  group_by(month) |>
  summarize(Dark = sum(volumeoffmktcross + volumeonmktcross +
                         volumecentrept - volumecentreptcross, na.rm = TRUE),
          Lit = sum(volumeonmkt + volumeoffmkt -
                      (volumeoffmktcross + volumeonmktcross +
                        volumecentrept - volumecentreptcross), na.rm = TRUE)) |>
  pivot_longer(cols = -month, names_to = "market", values_to = "volume") |>
  ggplot(aes(x = month, y = volume, color = market)) +
  geom_line() +
```

```
theme(legend.position = "top")
```



Figure 10: Trading activity: Dark versus lit market

The ASX Centre Point matching system provides a market for dark pool liquidity. As such, Centre Point trades are a subset of on-market trades. More information on ASX Centre Point can be found on the ASX website. The composition of each market segment is displayed at the top of the si_au_prc_daily tab in our data dictionary for this service. Figure 11 shows the average proportion of on-market non-crossing trades that are directed through ASX Centre Point over time.

```
si_au_prc_daily |>
  filter(valueonmkt > 0, valuecentrept > 0) |>
  mutate(month = floor_date(date, "month")) |>
  group_by(month) |>
  summarize(AvgPropCentrePtNonCross =
            mean((valuecentrept - valuecentreptcross) /
                  (valueonmkt - valueonmktcross), na.rm = TRUE)) |>
  ggplot(aes(x = month, y = AvgPropCentrePtNonCross)) +
  geom_line()
```

```
Warning: Removed 7 rows containing missing values or values outside the scale range
(`geom_line()`).
```

Figure 11: Centre Point market share

It is a simple matter to focus on particular market segments. For example, the previous query can be targeted on companies whose market capitalisation is less than $50 million, with results depicted in Figure 12.

```
si_au_prc_daily |>
  filter(close * shares < 50000000,
         valueonmkt > 0, valuecentrept > 0) |>
  mutate(month = floor_date(date, "month")) |>
  group_by(month) |>
  summarize(AvgPropCentrePtNonCross =
              mean((valuecentrept - valuecentreptcross) /
                     (valueonmkt - valueonmktcross), na.rm = TRUE)) |>
  ggplot(aes(x = month, y = AvgPropCentrePtNonCross)) +
  geom_line()
```

Warning: Removed 5 rows containing missing values or values outside the scale range
(`geom_line()`).

Figure 12: Centre Point market share: Small caps

Finally, Figure 13 shows the share of the dark market volumes traded on Centre Point.

```
si_au_prc_daily |>
  filter(valueonmkt > 0, valuecentrept > 0) |>
  mutate(Dark = volumeoffmktcross + volumeonmktcross +
                        volumecentrept - volumecentreptcross,
          CentrePt = volumecentrept - volumecentreptcross) |>
  mutate(month = floor_date(date, "month")) |>
  group_by(month) |>
  summarize(AvgPropCentrePtDark = mean(CentrePt / Dark, na.rm = TRUE)) |>
  ggplot(aes(x = month, y = AvgPropCentrePtDark)) +
  geom_line()
```

Figure 13: Centre Point share of dark market

```
dbDisconnect(db)
```